

Lessons Learned

Security at an Open Source Startup

Chris Sandulow @jfalken and **Andreas Nilsson** @agralius



\$ whois []

This Talk

InfoSec at a Startup

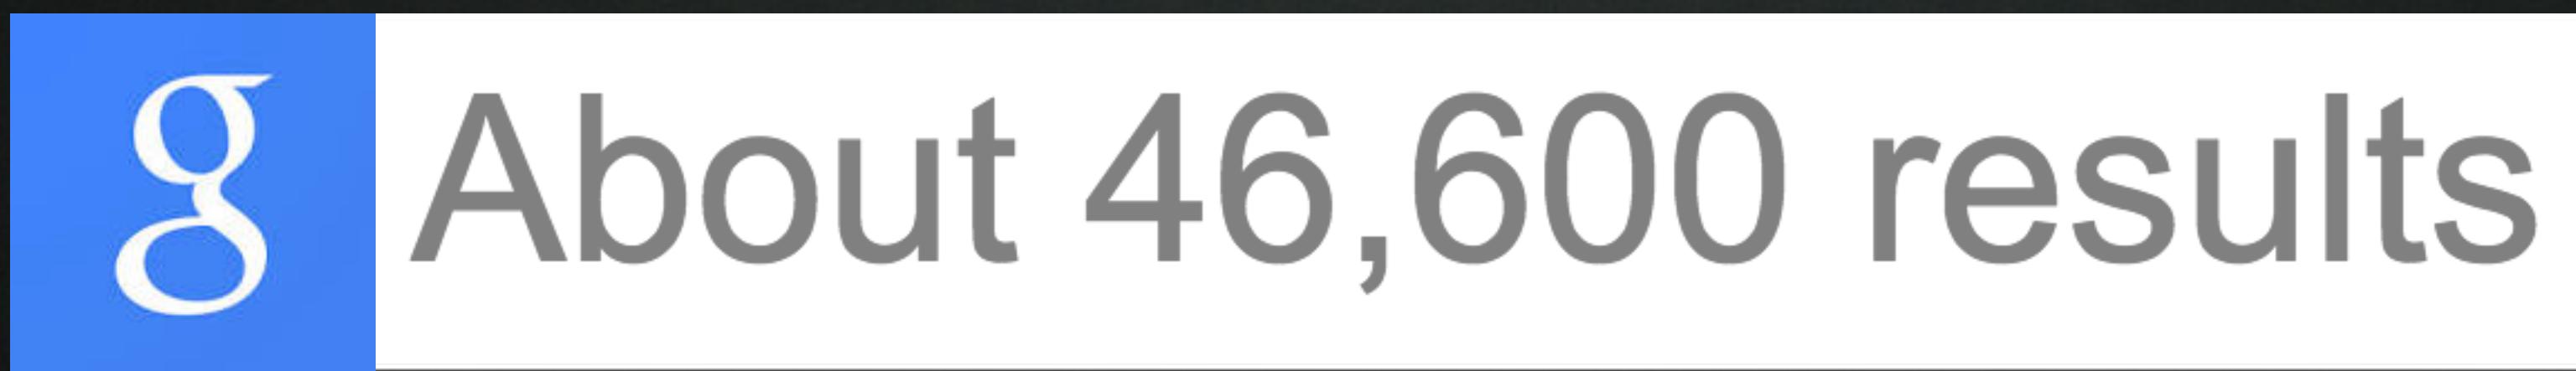
Organizational / Enterprise Security

Application Security / SDL

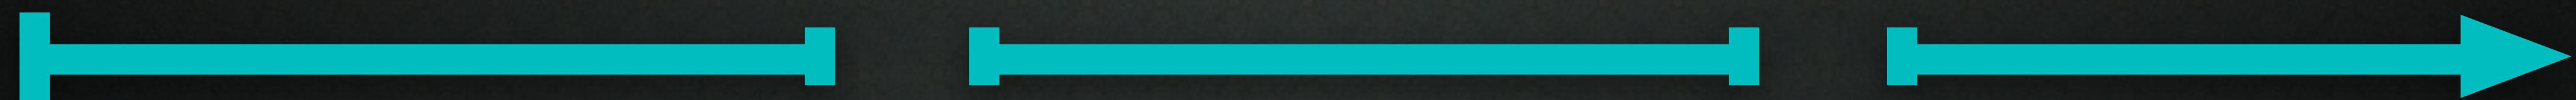
Tools

**“We take security
very seriously”**

**“We take security
very seriously”**



Size Matters



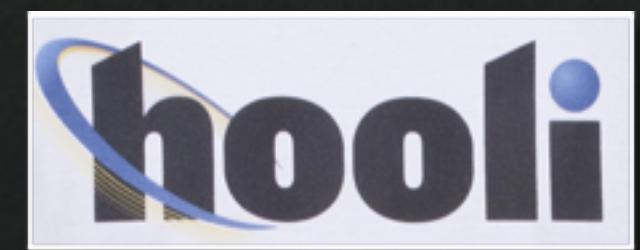
0 to mid-size-ish



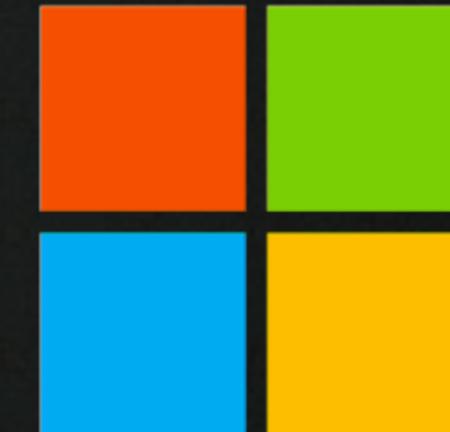
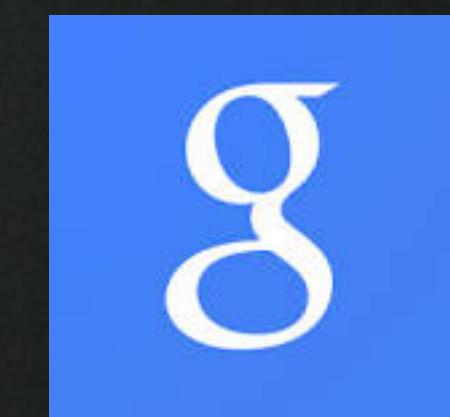
Mid-Cap / Large-Cap

Banks

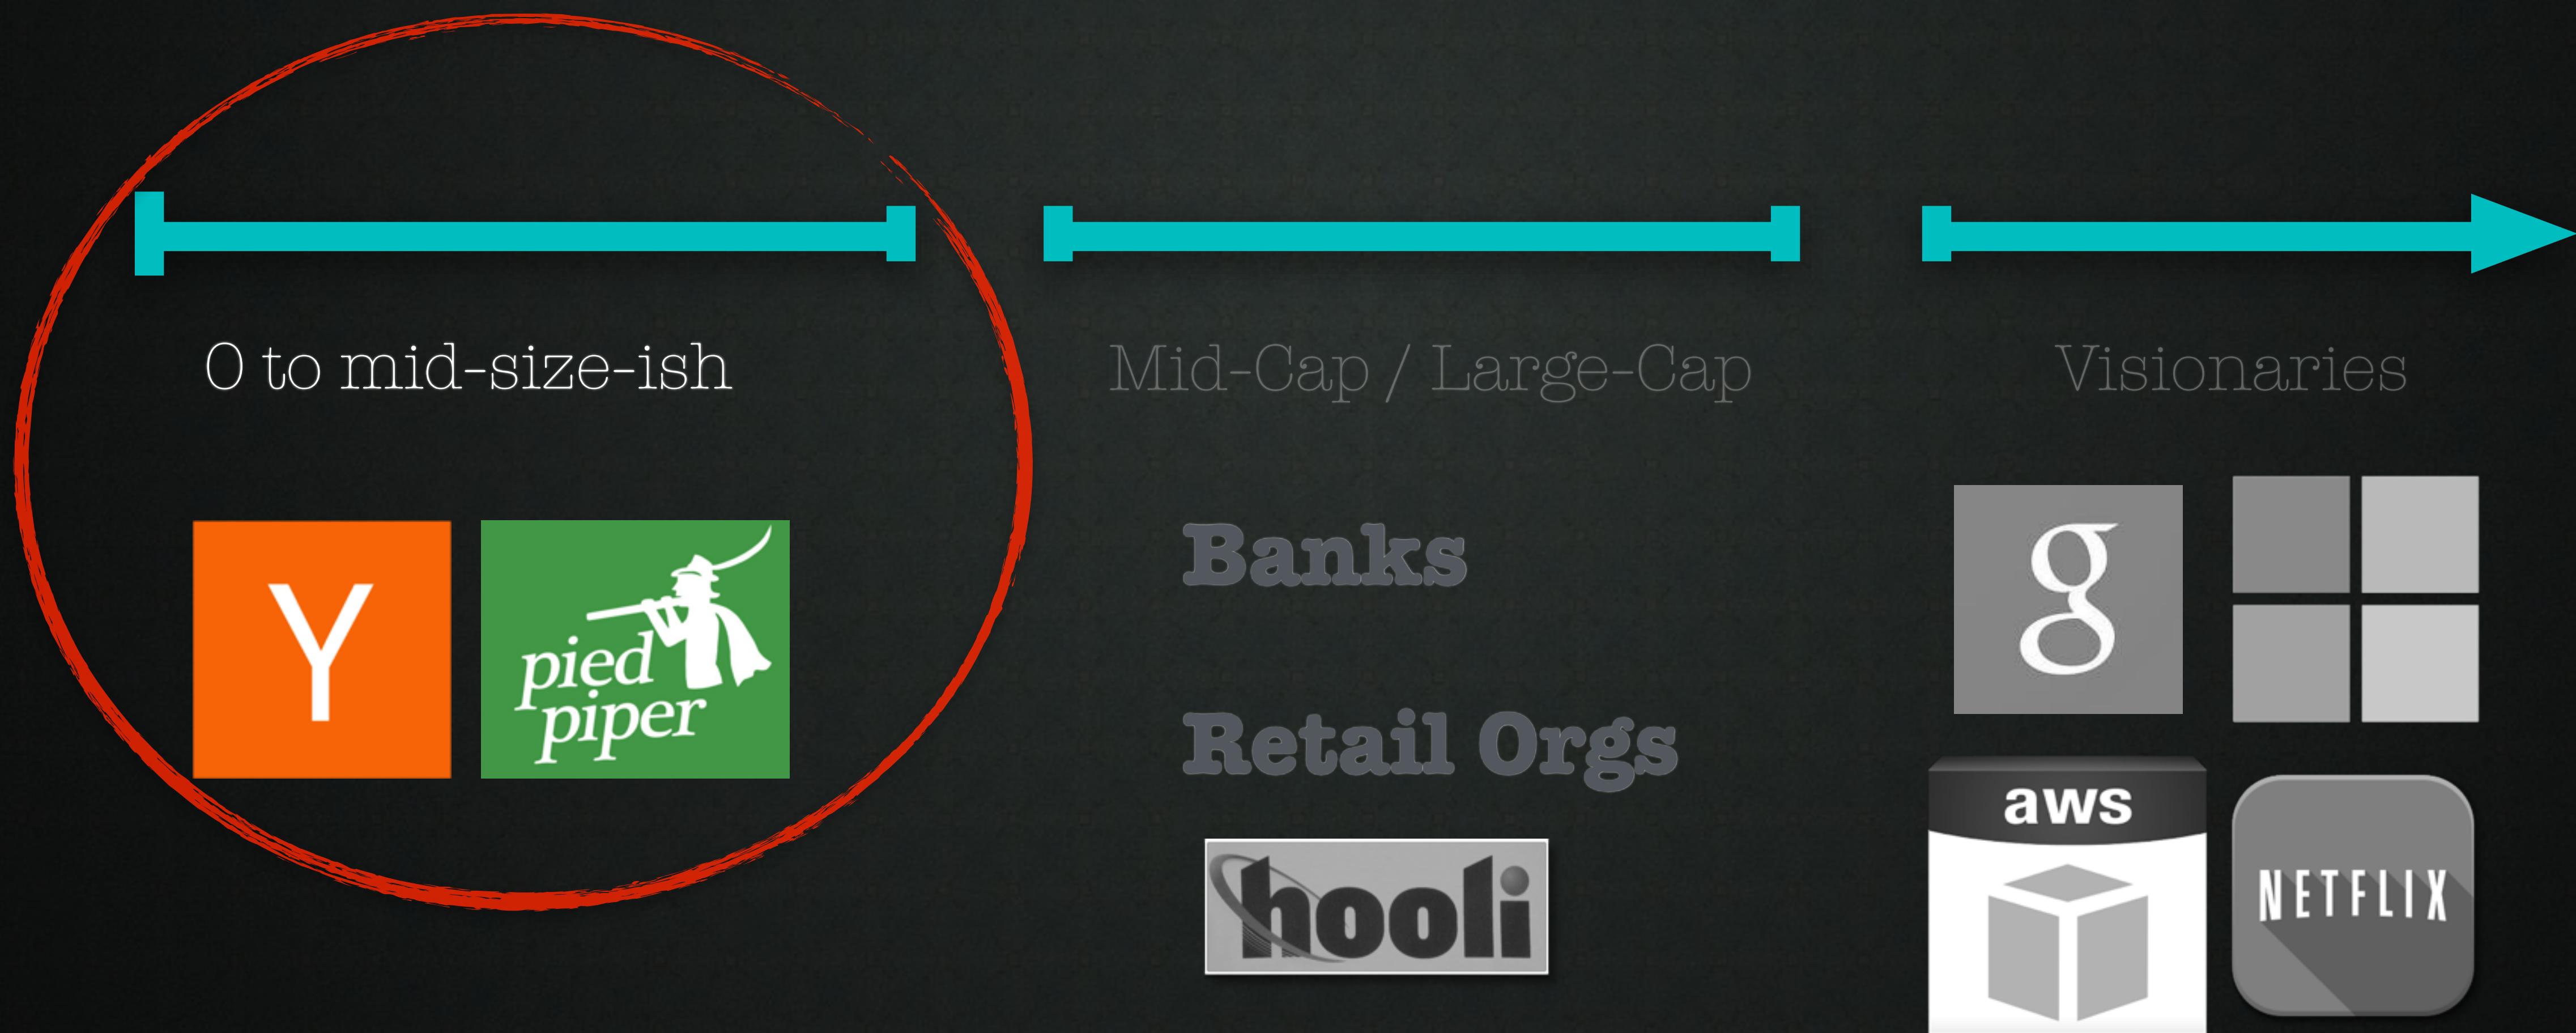
Retail Orgs



Visionaries



Size Matters



People



**“I can’t recommend anyone to be
a generalist... I don’t think it’s
possible ... to
spectrum generalist.”**

Dan Geer

<http://geertinho.net/geer.sourceboston.txt>



When to build a Security Function?

- ❖ The Sooner the Better, duh!
- ❖ B2C, sooner than B2B; Regulated sooner than Non
- ❖ Don't let compliance solely drive the program
- ❖ **You probably need atleast 2, not just 1 person.**
- ❖ **Once you're shipping a product/service, you should have a dedicated security person**

Who? and How Many?

- ❖ Purple Unicorn; field diversification

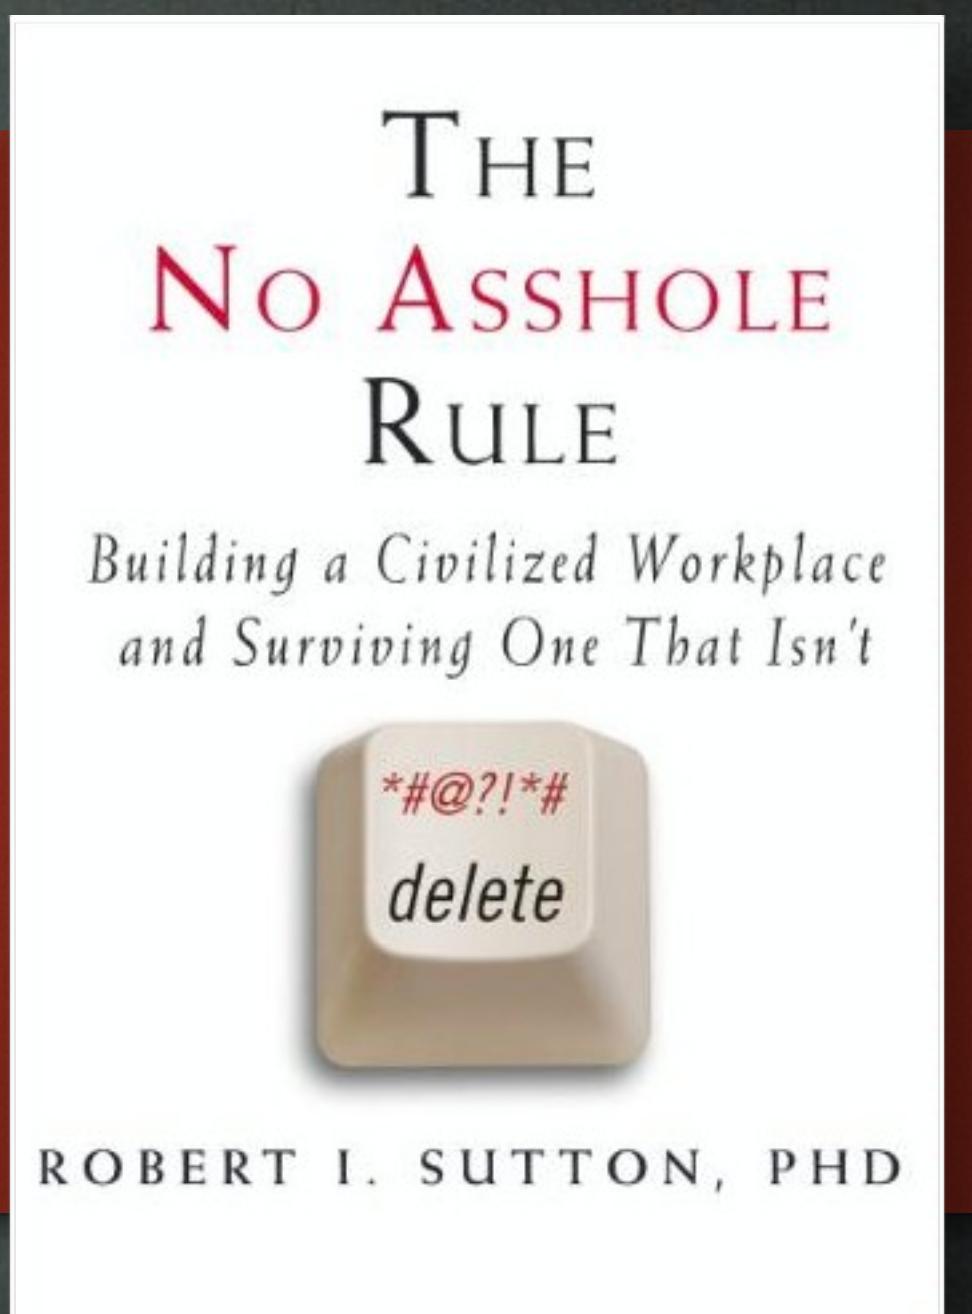
“...there is negative unemployment for highly-skilled security professionals.”

— Lee Kushner

- ❖ Skills: “Code to an API”
- ❖ Capabilities: Builders vs Breakers
- ❖ Sizing and Scaling: 1% - 2% of Org *



@iodboi



Org / Enterprise Sec



First Step: InfoSec Program and Baseline

- ❖ Baseline program goal; vision; what to achieve
- ❖ Driving Adoption and FUD
- ❖ Measuring Progress
- ❖ Transparency
- ❖ ***Attack Driven Defense*** – Zane Lackey



@zanelackey

All the A's

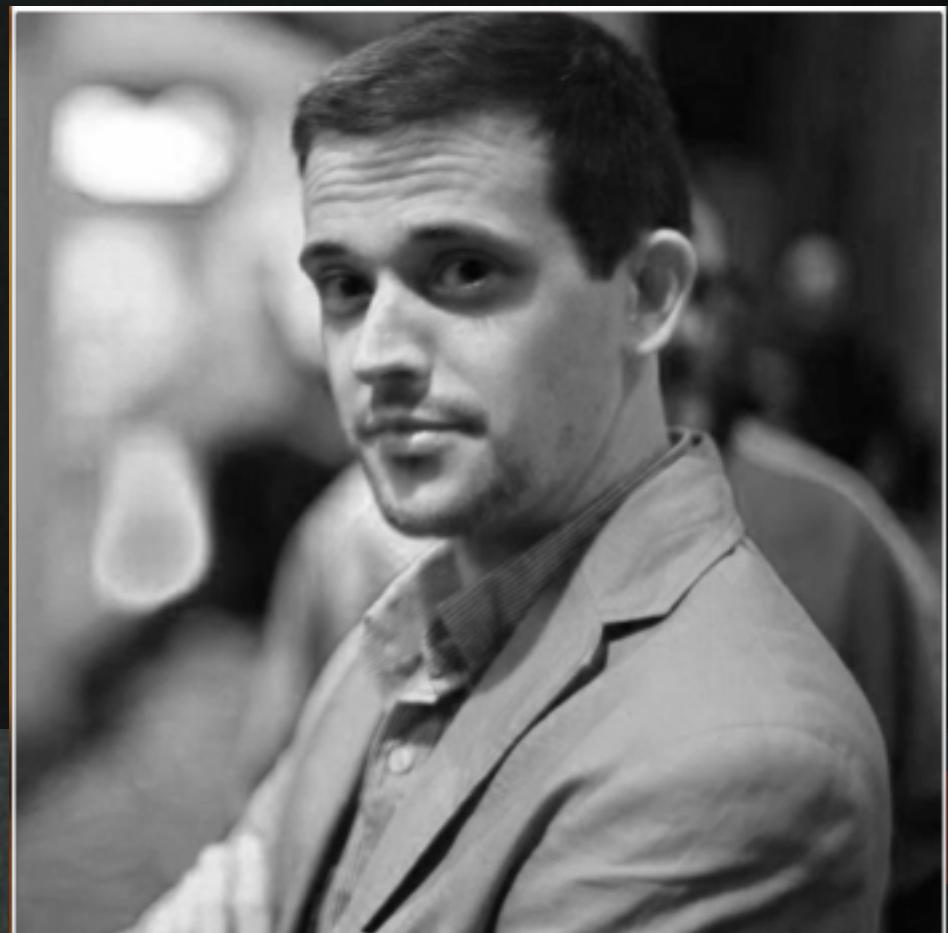
0000000:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000010:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000020:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000030:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000040:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000050:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000060:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000070:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000080:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000090:	4141	4141	4141	4141	4141	4141	4141	4141	4141
00000a0:	4141	4141	4141	4141	4141	4141	4141	4141	4141
00000b0:	4141	4141	4141	4141	4141	4141	4141	4141	4141
00000c0:	4141	4141	4141	4141	4141	4141	4141	4141	4141
00000d0:	4141	4141	4141	4141	4141	4141	4141	4141	4141
00000e0:	4141	4141	4141	4141	4141	4141	4141	4141	4141
00000f0:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000100:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000110:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000120:	4141	4141	4141	4141	4141	4141	4141	4141	4141
0000130:	4141	4141	4141	4141	4141	4141	4141	4141	4141

**AuthN
Accounts
Assets
Auditing**

“...only two effective
security technologies,
firewalls and 2fa”

@chrisrohlf

<https://twitter.com/chrisrohlf/status/574639358101352449>



Curtailing Identities

curtail – kər'tāl, verb

Reduce in extent or quantity;

Impose a restriction on.

Accounts: Workforce Lifecycle Management

- ❖ Onboarding, Off-boarding, Role-Change(s)
- ❖ Reduce Identities; Centralize AuthN and AuthZ
- ❖ Federated Identity (SAML)
- ❖ Security Control reduces risk and is more convenient
- ❖ HR and IT; Documentation
- ❖ Automate process audits

Assets

- ❖ Thinking about Assets
- ❖ Cloud vs. Meatspace
- ❖ Servers vs Laptops/Desktop
- ❖ Tagging and Schema
- ❖ Diffs

Asset Monitoring

Check the Deltas

- ❖ Changes in key infra (e.g., load balancers, ACLs, admin memberships, etc.)
- ❖ Use Version Control tools
- ❖ Develop tooling to automate this reporting
- ❖ Stimulus-response cycle



Vulnerability Scanning

- ❖ Once assets are identified, we can scan them
- ❖ Don't expect to find much? Expect to be surprised
- ❖ Triage; don't patch all-things all-the-time
- ❖ Ownership, Patching, Procedures w/ DevOps team
- ❖ Reduce exposure, Reduce exposure, Reduce exposure

AppSec



Establish an SDL

- ❖ Secure Development Lifecycle (coined by Microsoft)
- ❖ Establish one that makes sense for your org.

Requirements

Implementation

Release

Design

Verification

Response

Training

SDL for a Startup

- ❖ Internal training of engineers
- ❖ Design reviews of all security features
- ❖ No commits without peer code review
- ❖ Use static code analysis
- ❖ Use third party for security assessment

Code Reviews

- ❖ Publish and enforce coding style guide
- ❖ No commits without review
- ❖ Use a common review tool
- ❖ In person reviews is the best training

Software Verification

- ❖ No new functionality without regression tests
- ❖ Automate and run tests in CI system
- ❖ Test unexpected and abnormal cases
- ❖ Use tools, ASAN (Address Sanitizer), Valgrind

Handling Vulnerabilities

- ❖ Tell people how to contact you for Security Issues
- ❖ Responsible Disclosure Policy
- ❖ PGP Key for encrypting sensitive comms
- ❖ Have a process to determine when to CVE an issue
- ❖ We've all seen tons of bad examples; don't be another. Be transparent and thankful.

Community / Bug Bounty

- ❖ Great way to scale a security program
- ❖ First time? Set Ceilings
- ❖ Be prepared for everything to be hit, firehose of emails
- ❖ Be Thankful, celebrate involvement

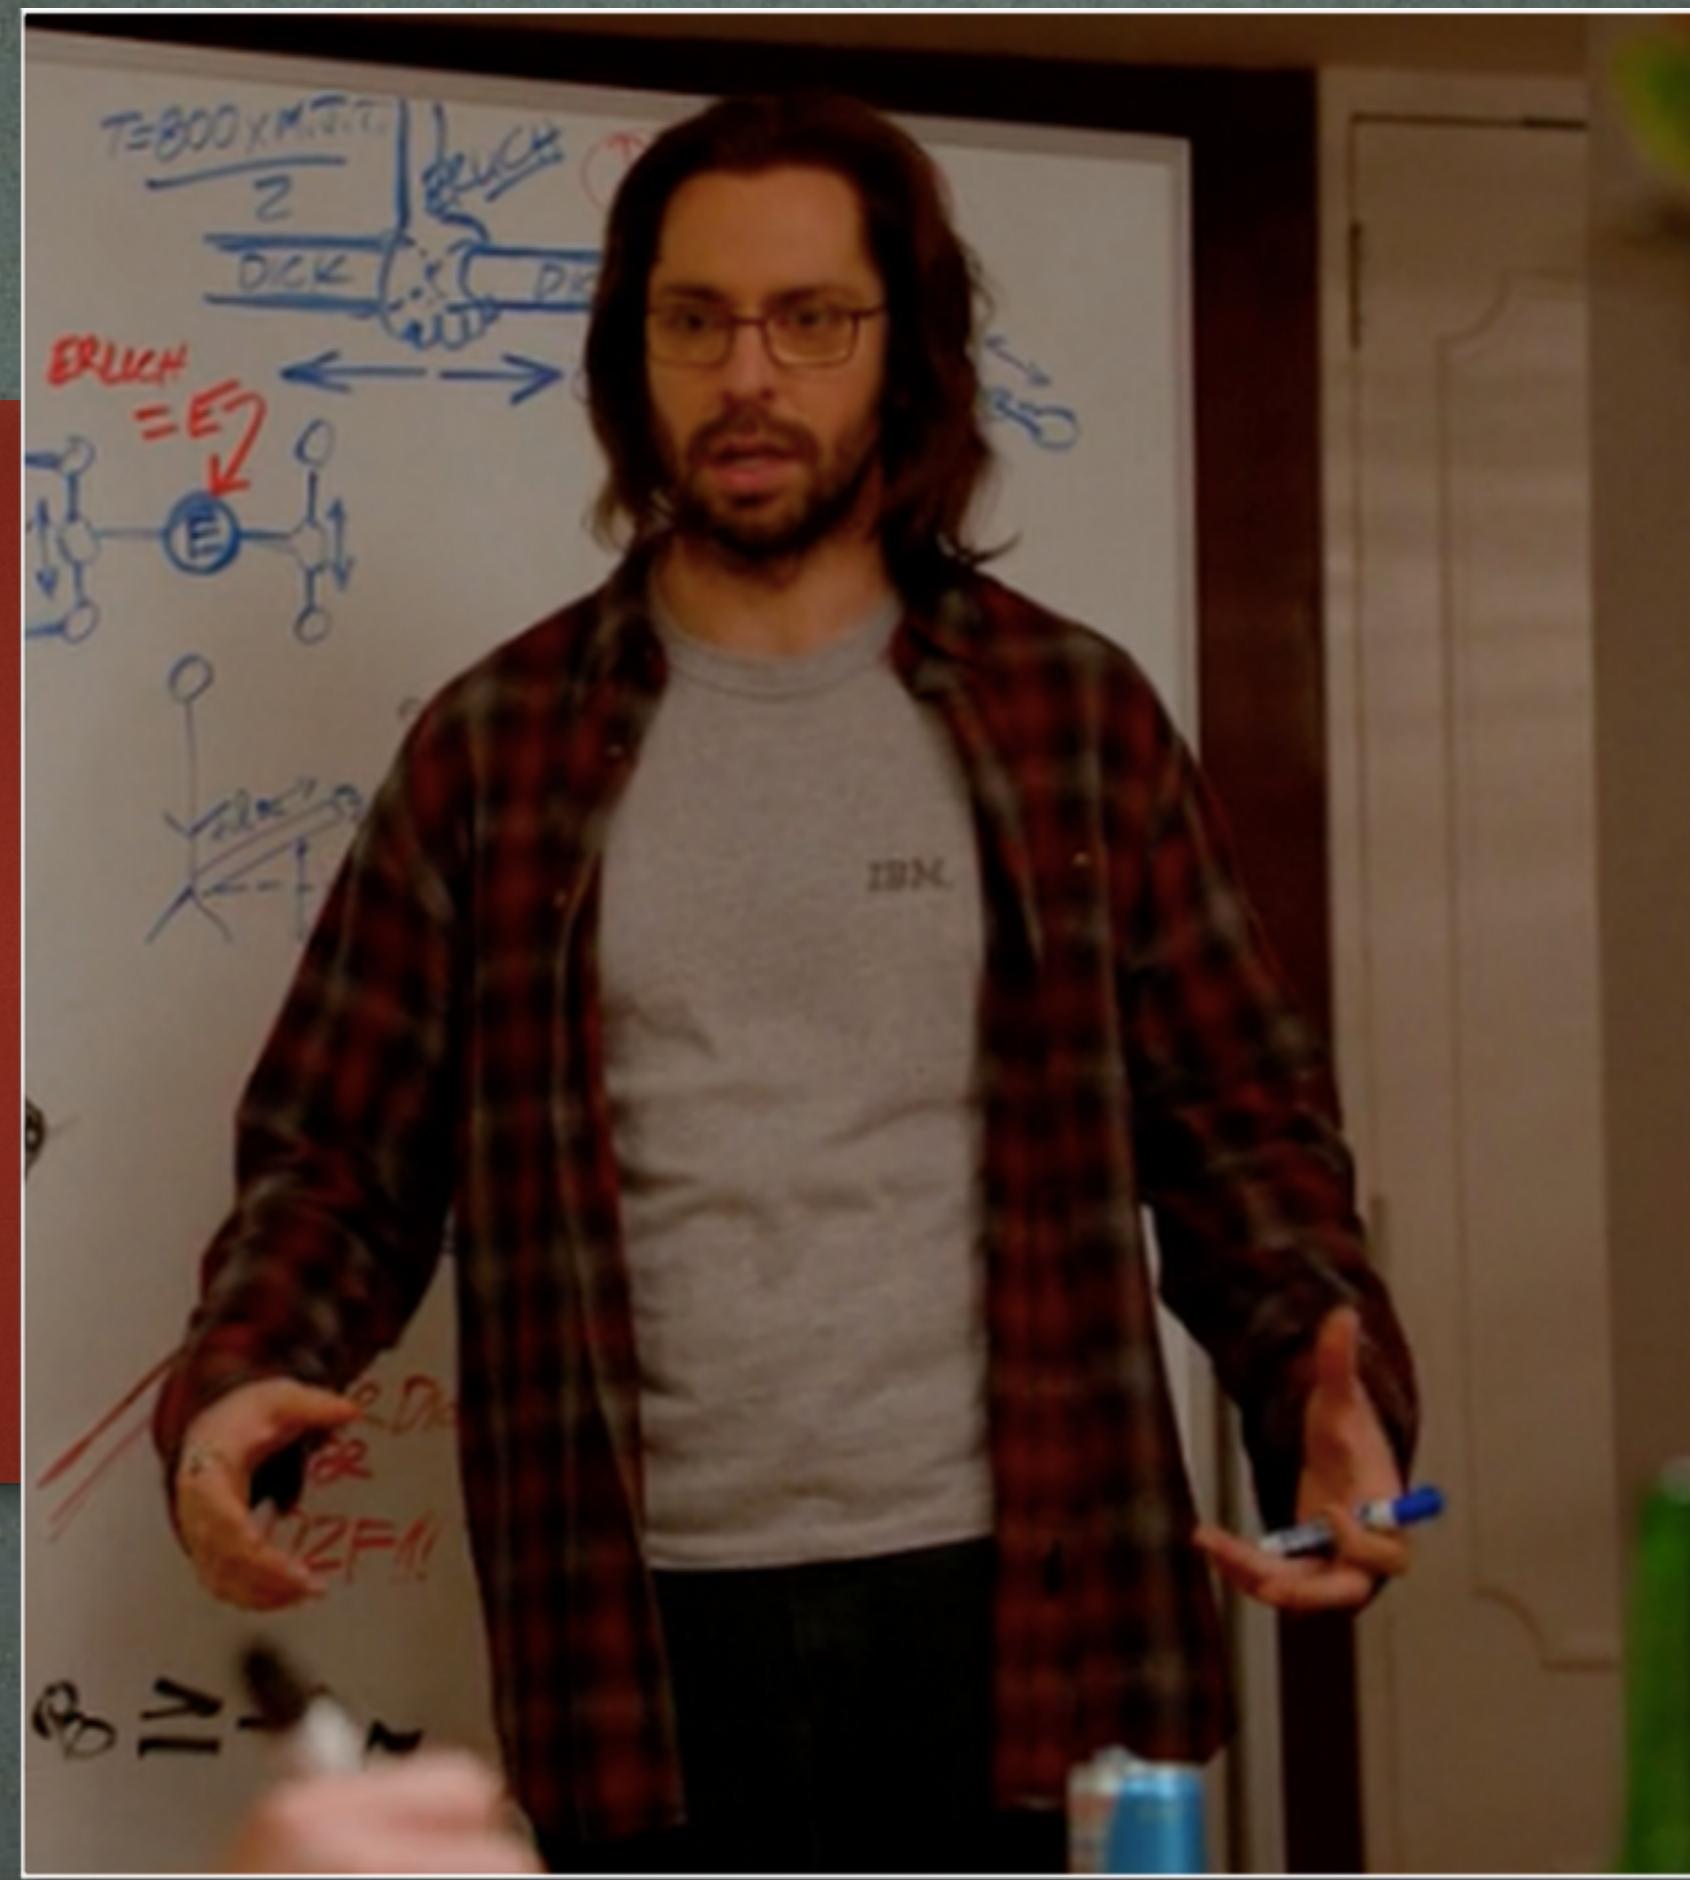
Secure and Sane Defaults

- ❖ For anything customer facing:
 - ❖ Secure Defaults save you grief later on
 - ❖ ‘Secure Defaults’ will hurt adoption as usability is affected; debate this
- ❖ Non-Customer Facing (ie, infrastructure, architecture, app-design)
 - ❖ You have no excuse to no make wise decisions from the get-go
 - ❖ Measure twice, cut once.

Open Sourcing

- ❖ Blessing
 - ❖ Get testing and reviews for free
 - ❖ Not possible to cheat through obfuscation
- ❖ and a curse
 - ❖ Coordinate security fixes in a public repo
 - ❖ Serve a user community and large customers concurrently

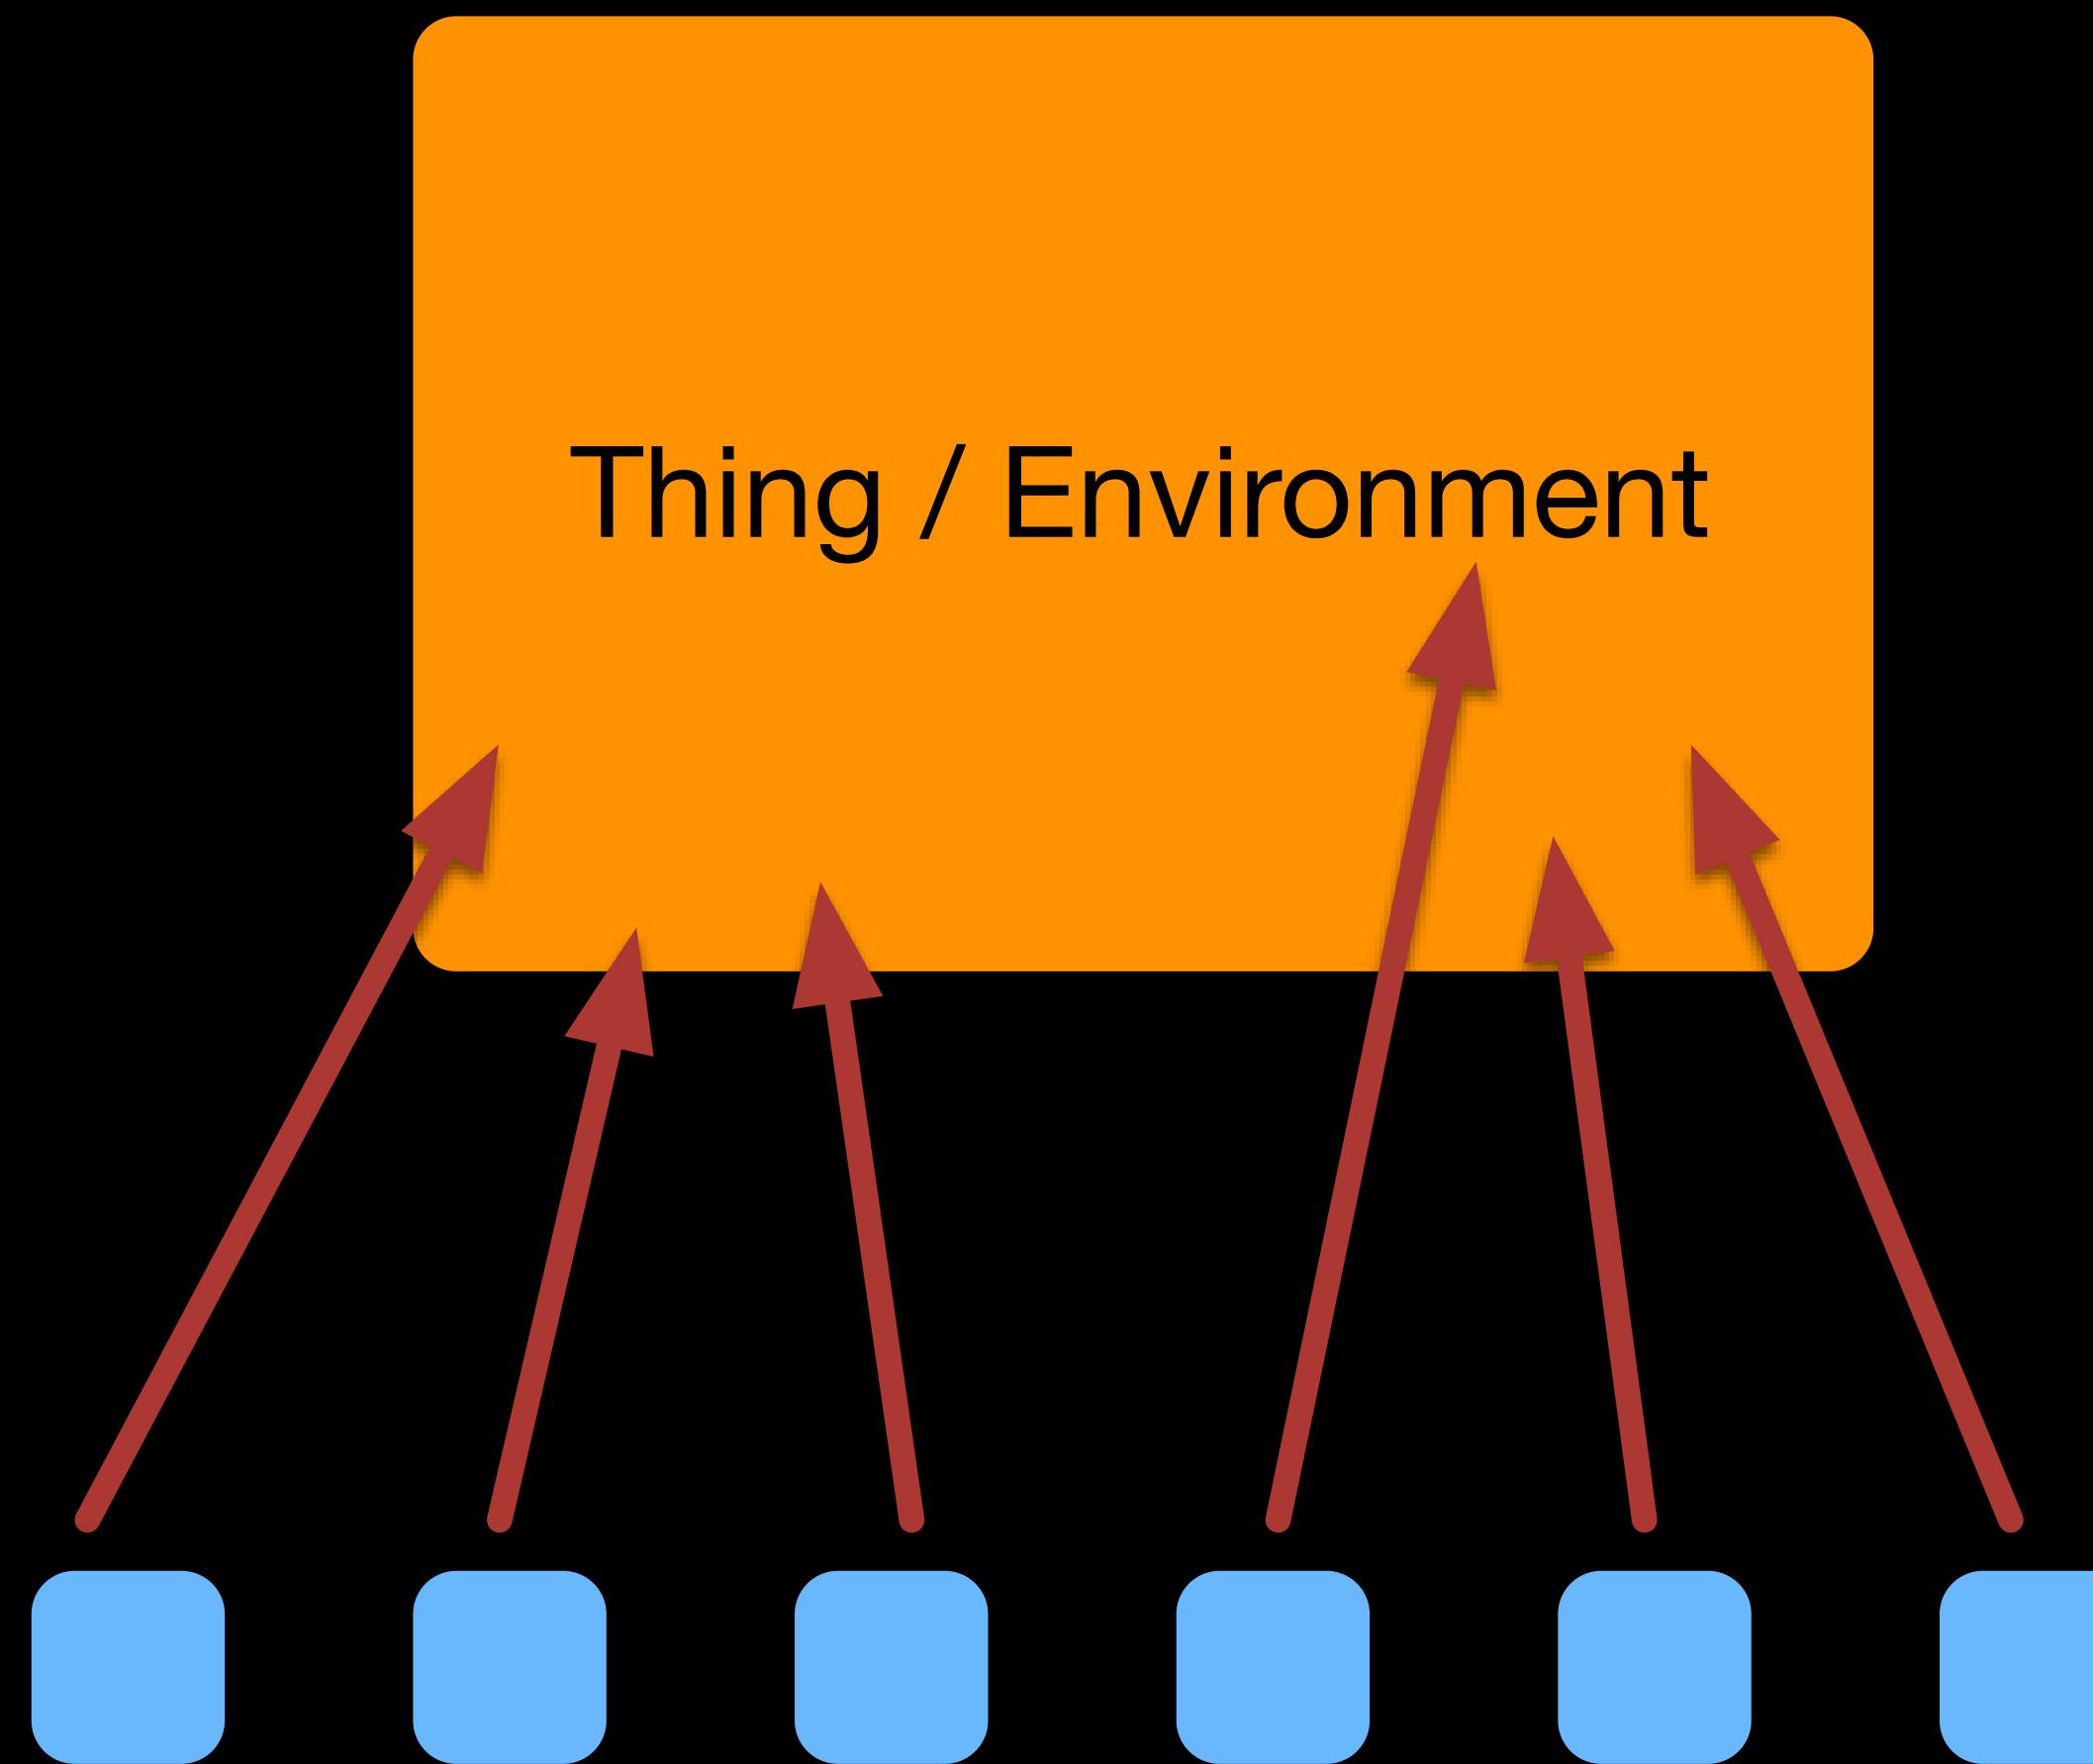
Architecting



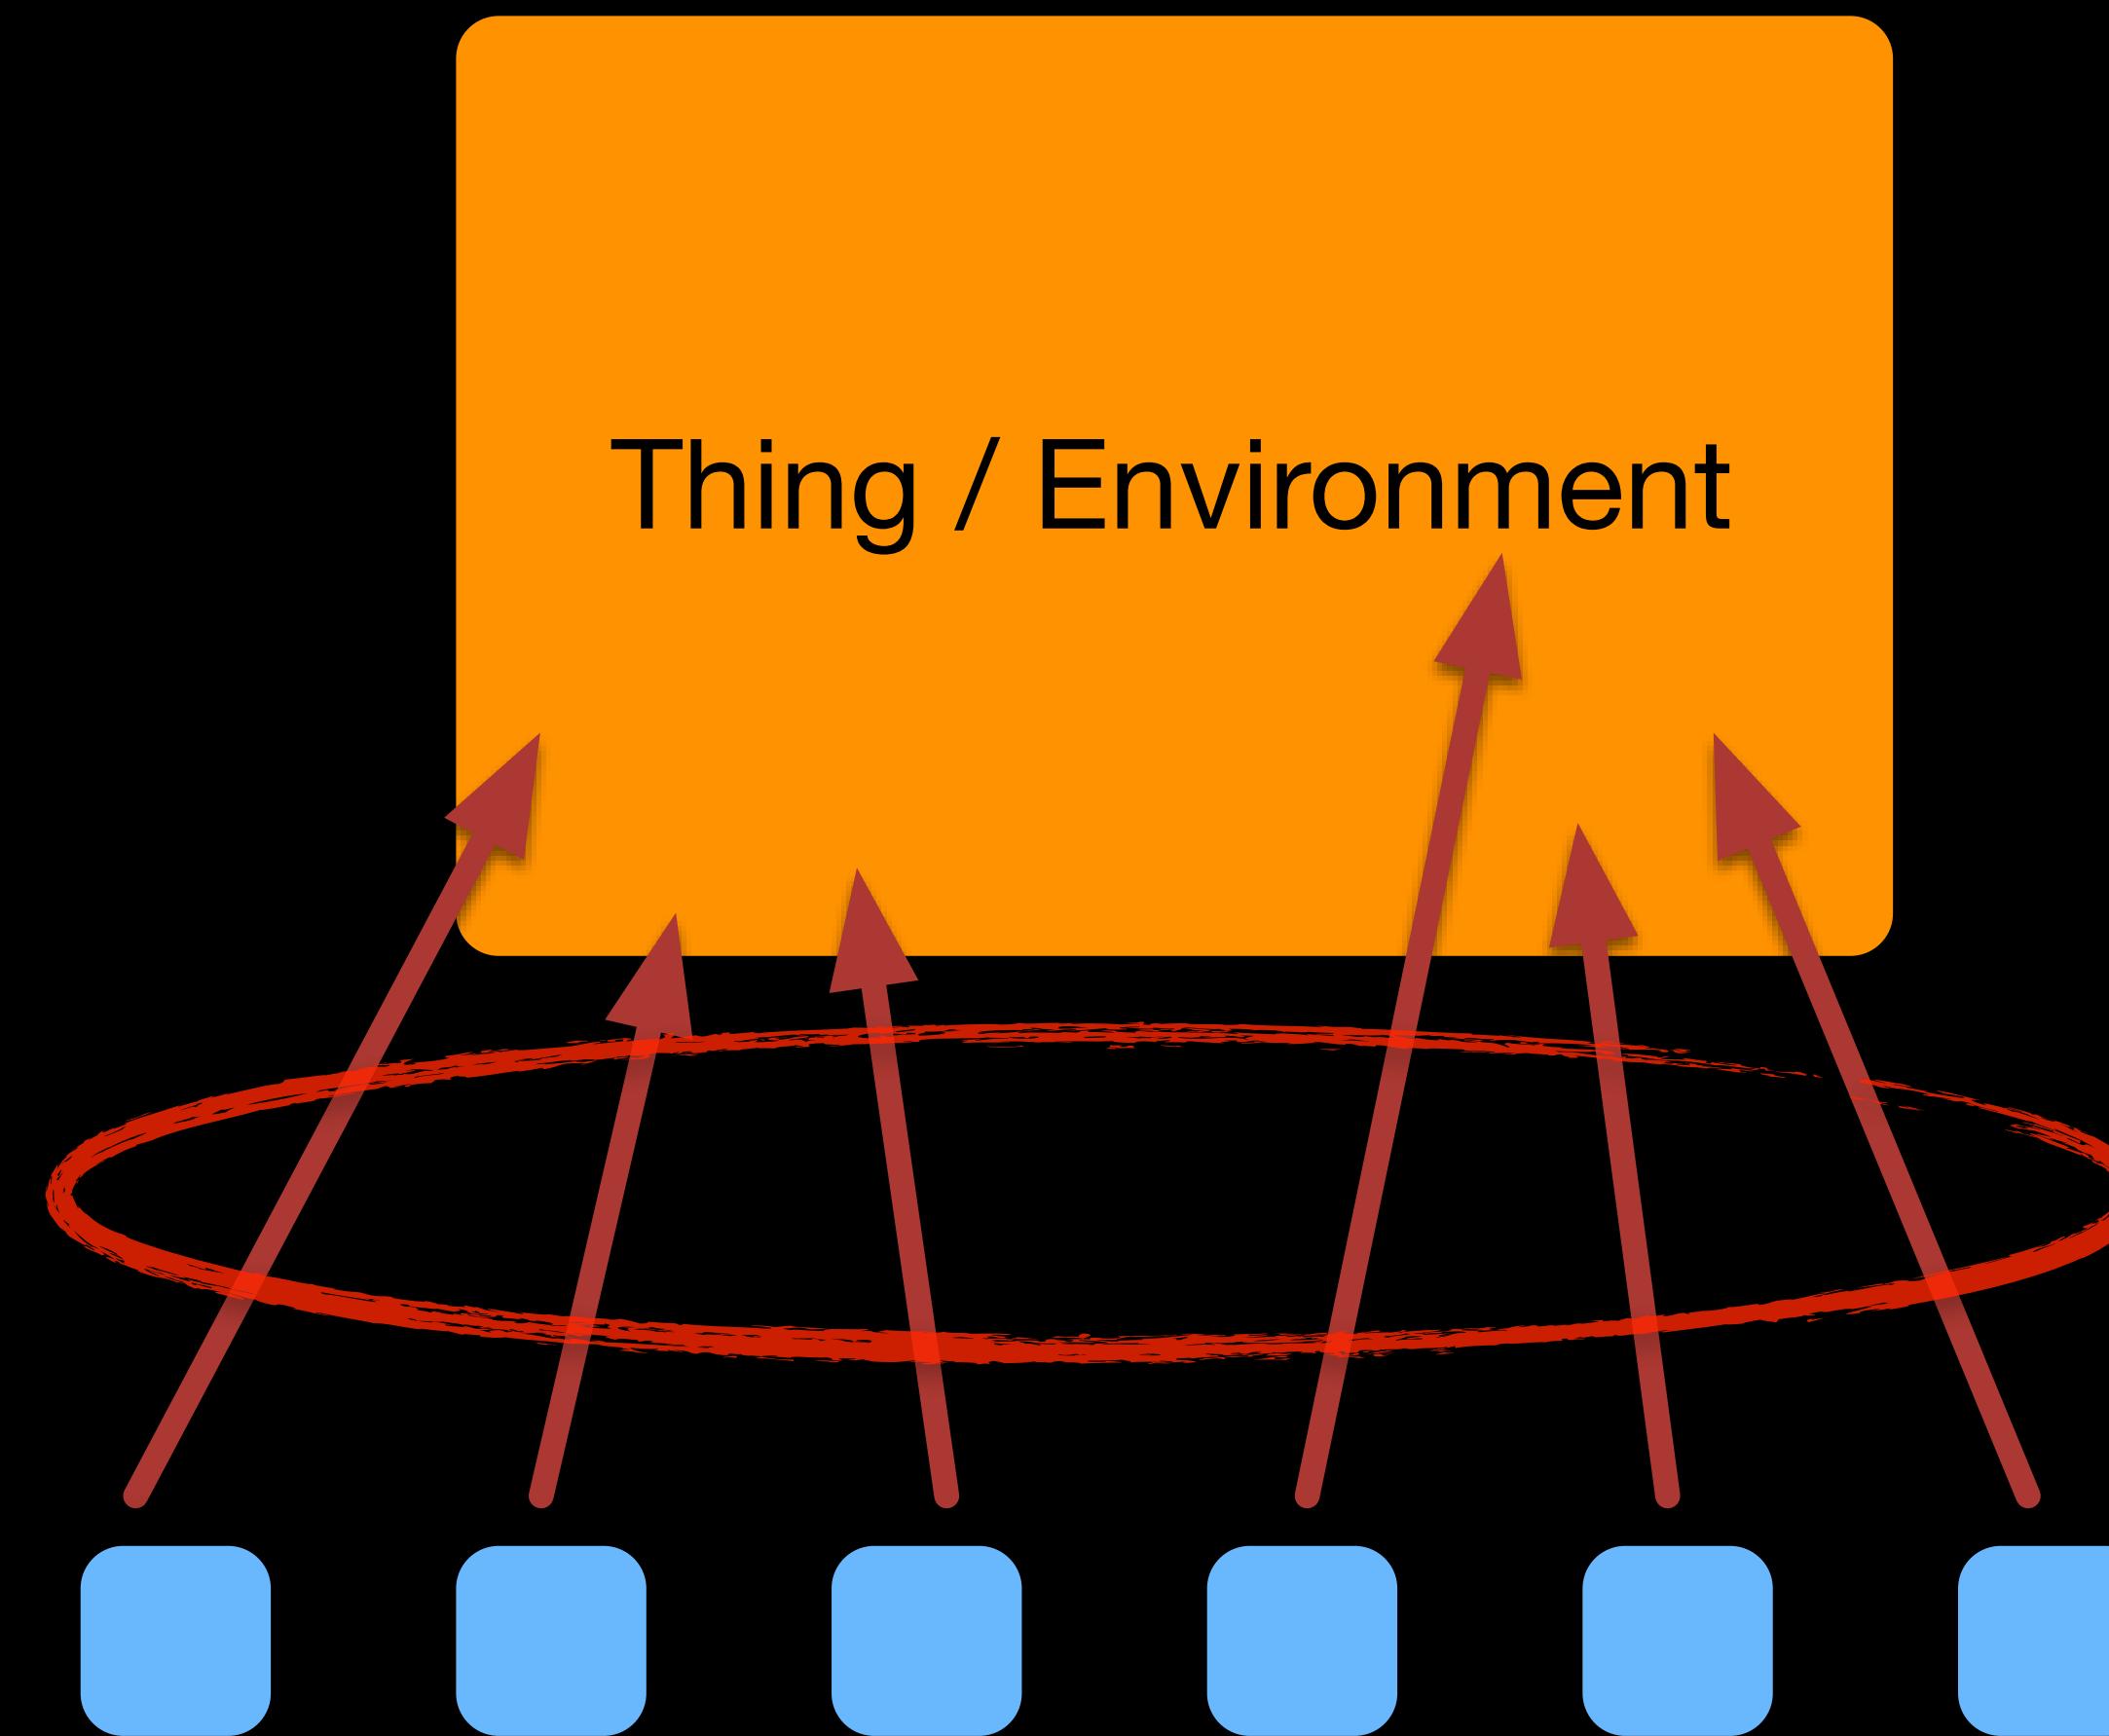
Note on Architectures

- ❖ Prefer loosely coupled applications
- ❖ Encourage Ephemeral Hosts
- ❖ Follow choke-point model
- ❖ Prefer auto-scaling / orchestration frameworks
- ❖ Find a way to not store secrets, or easily rotate them
 - ❖ e.g., ephemeral credentials

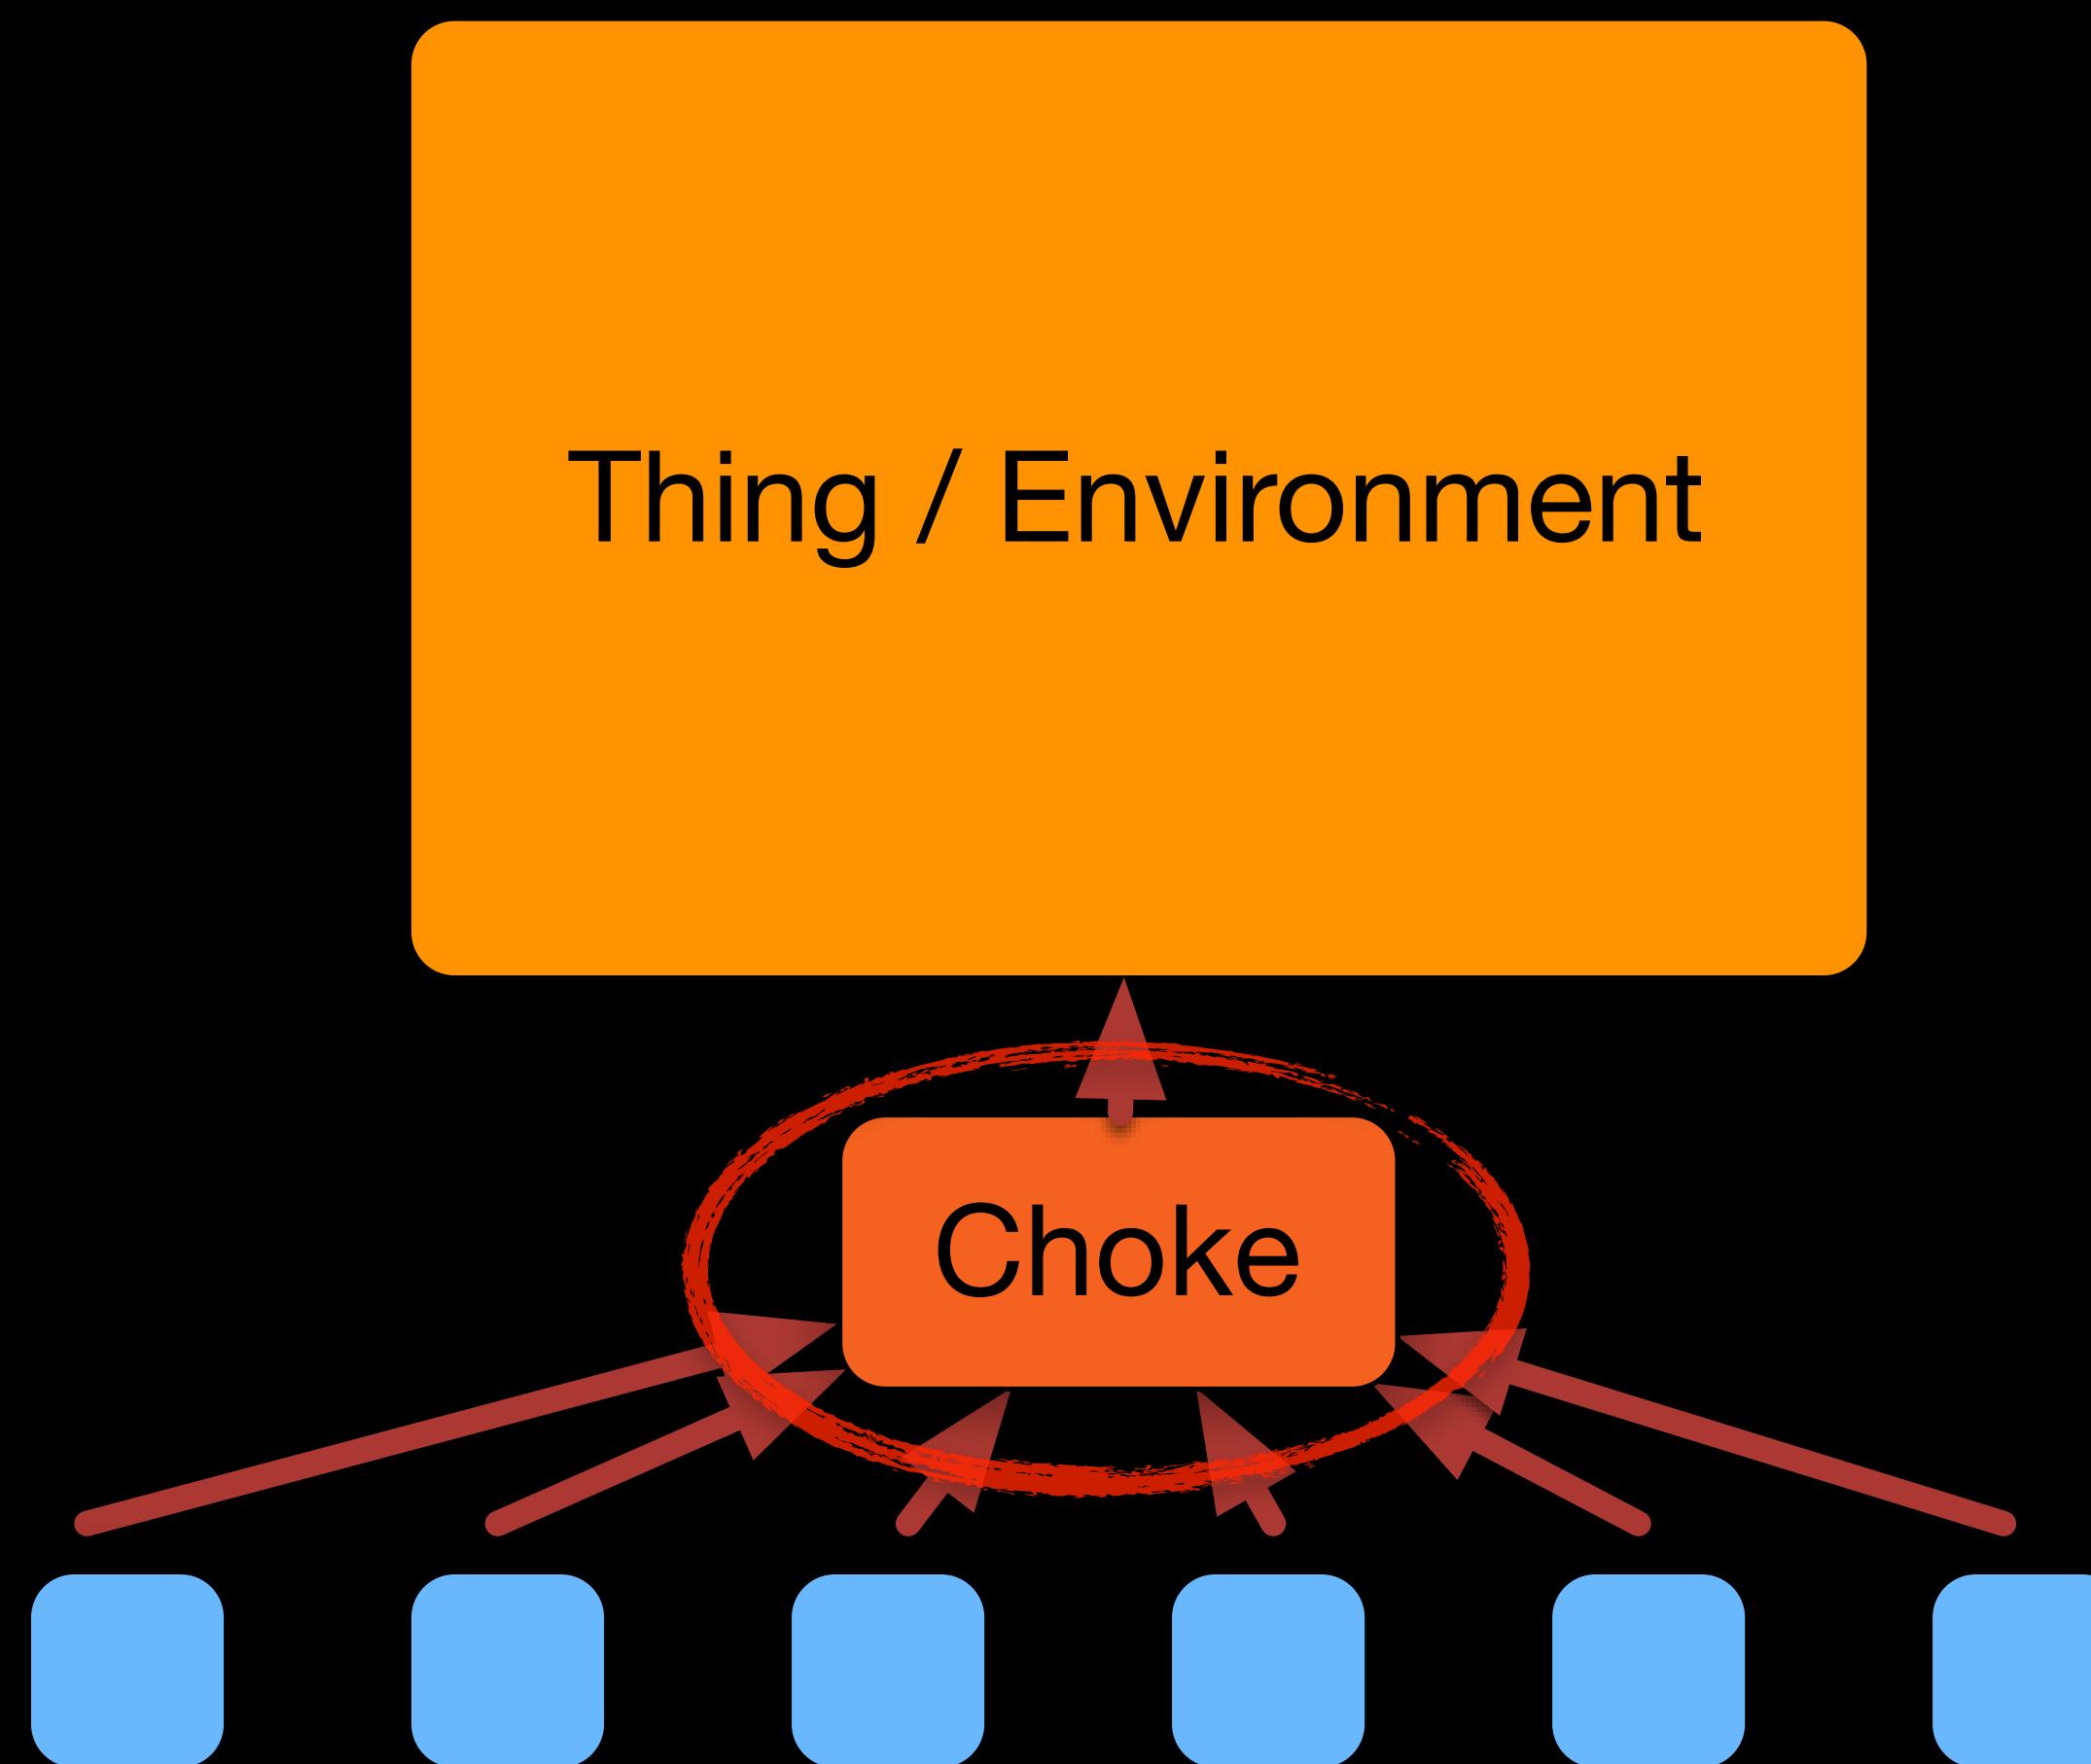
Architecting Choke Points



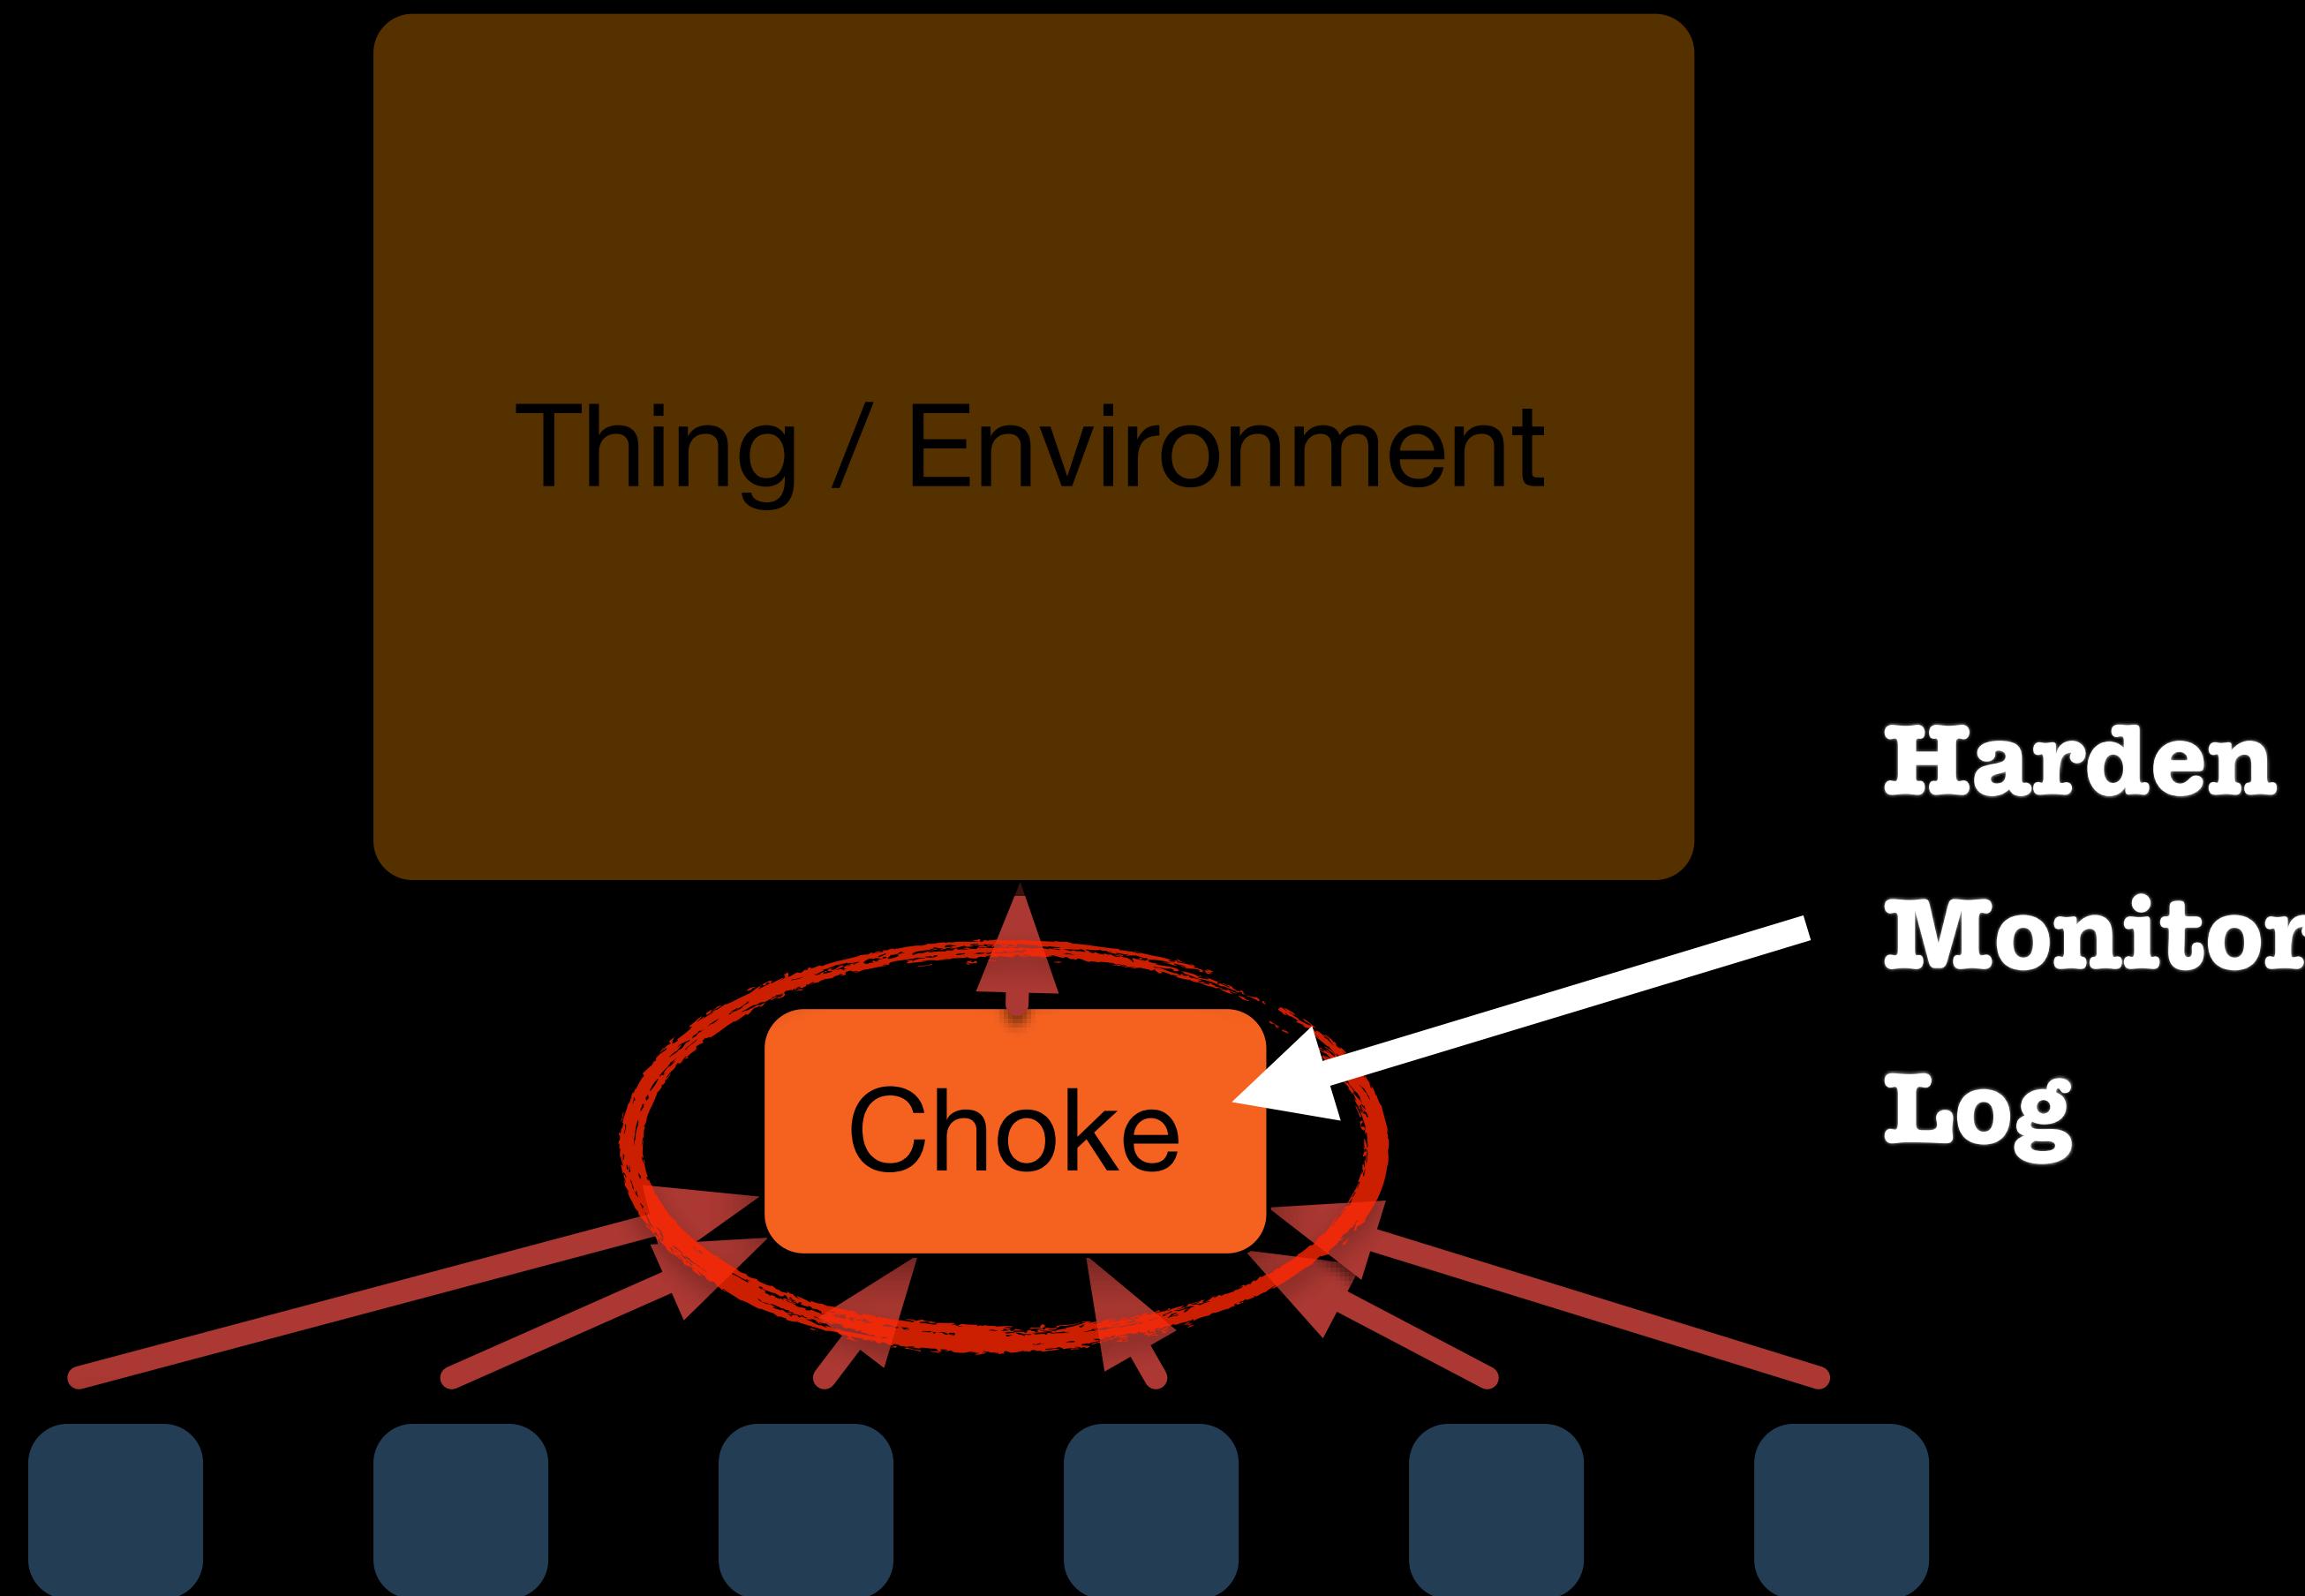
Architecting Choke Points



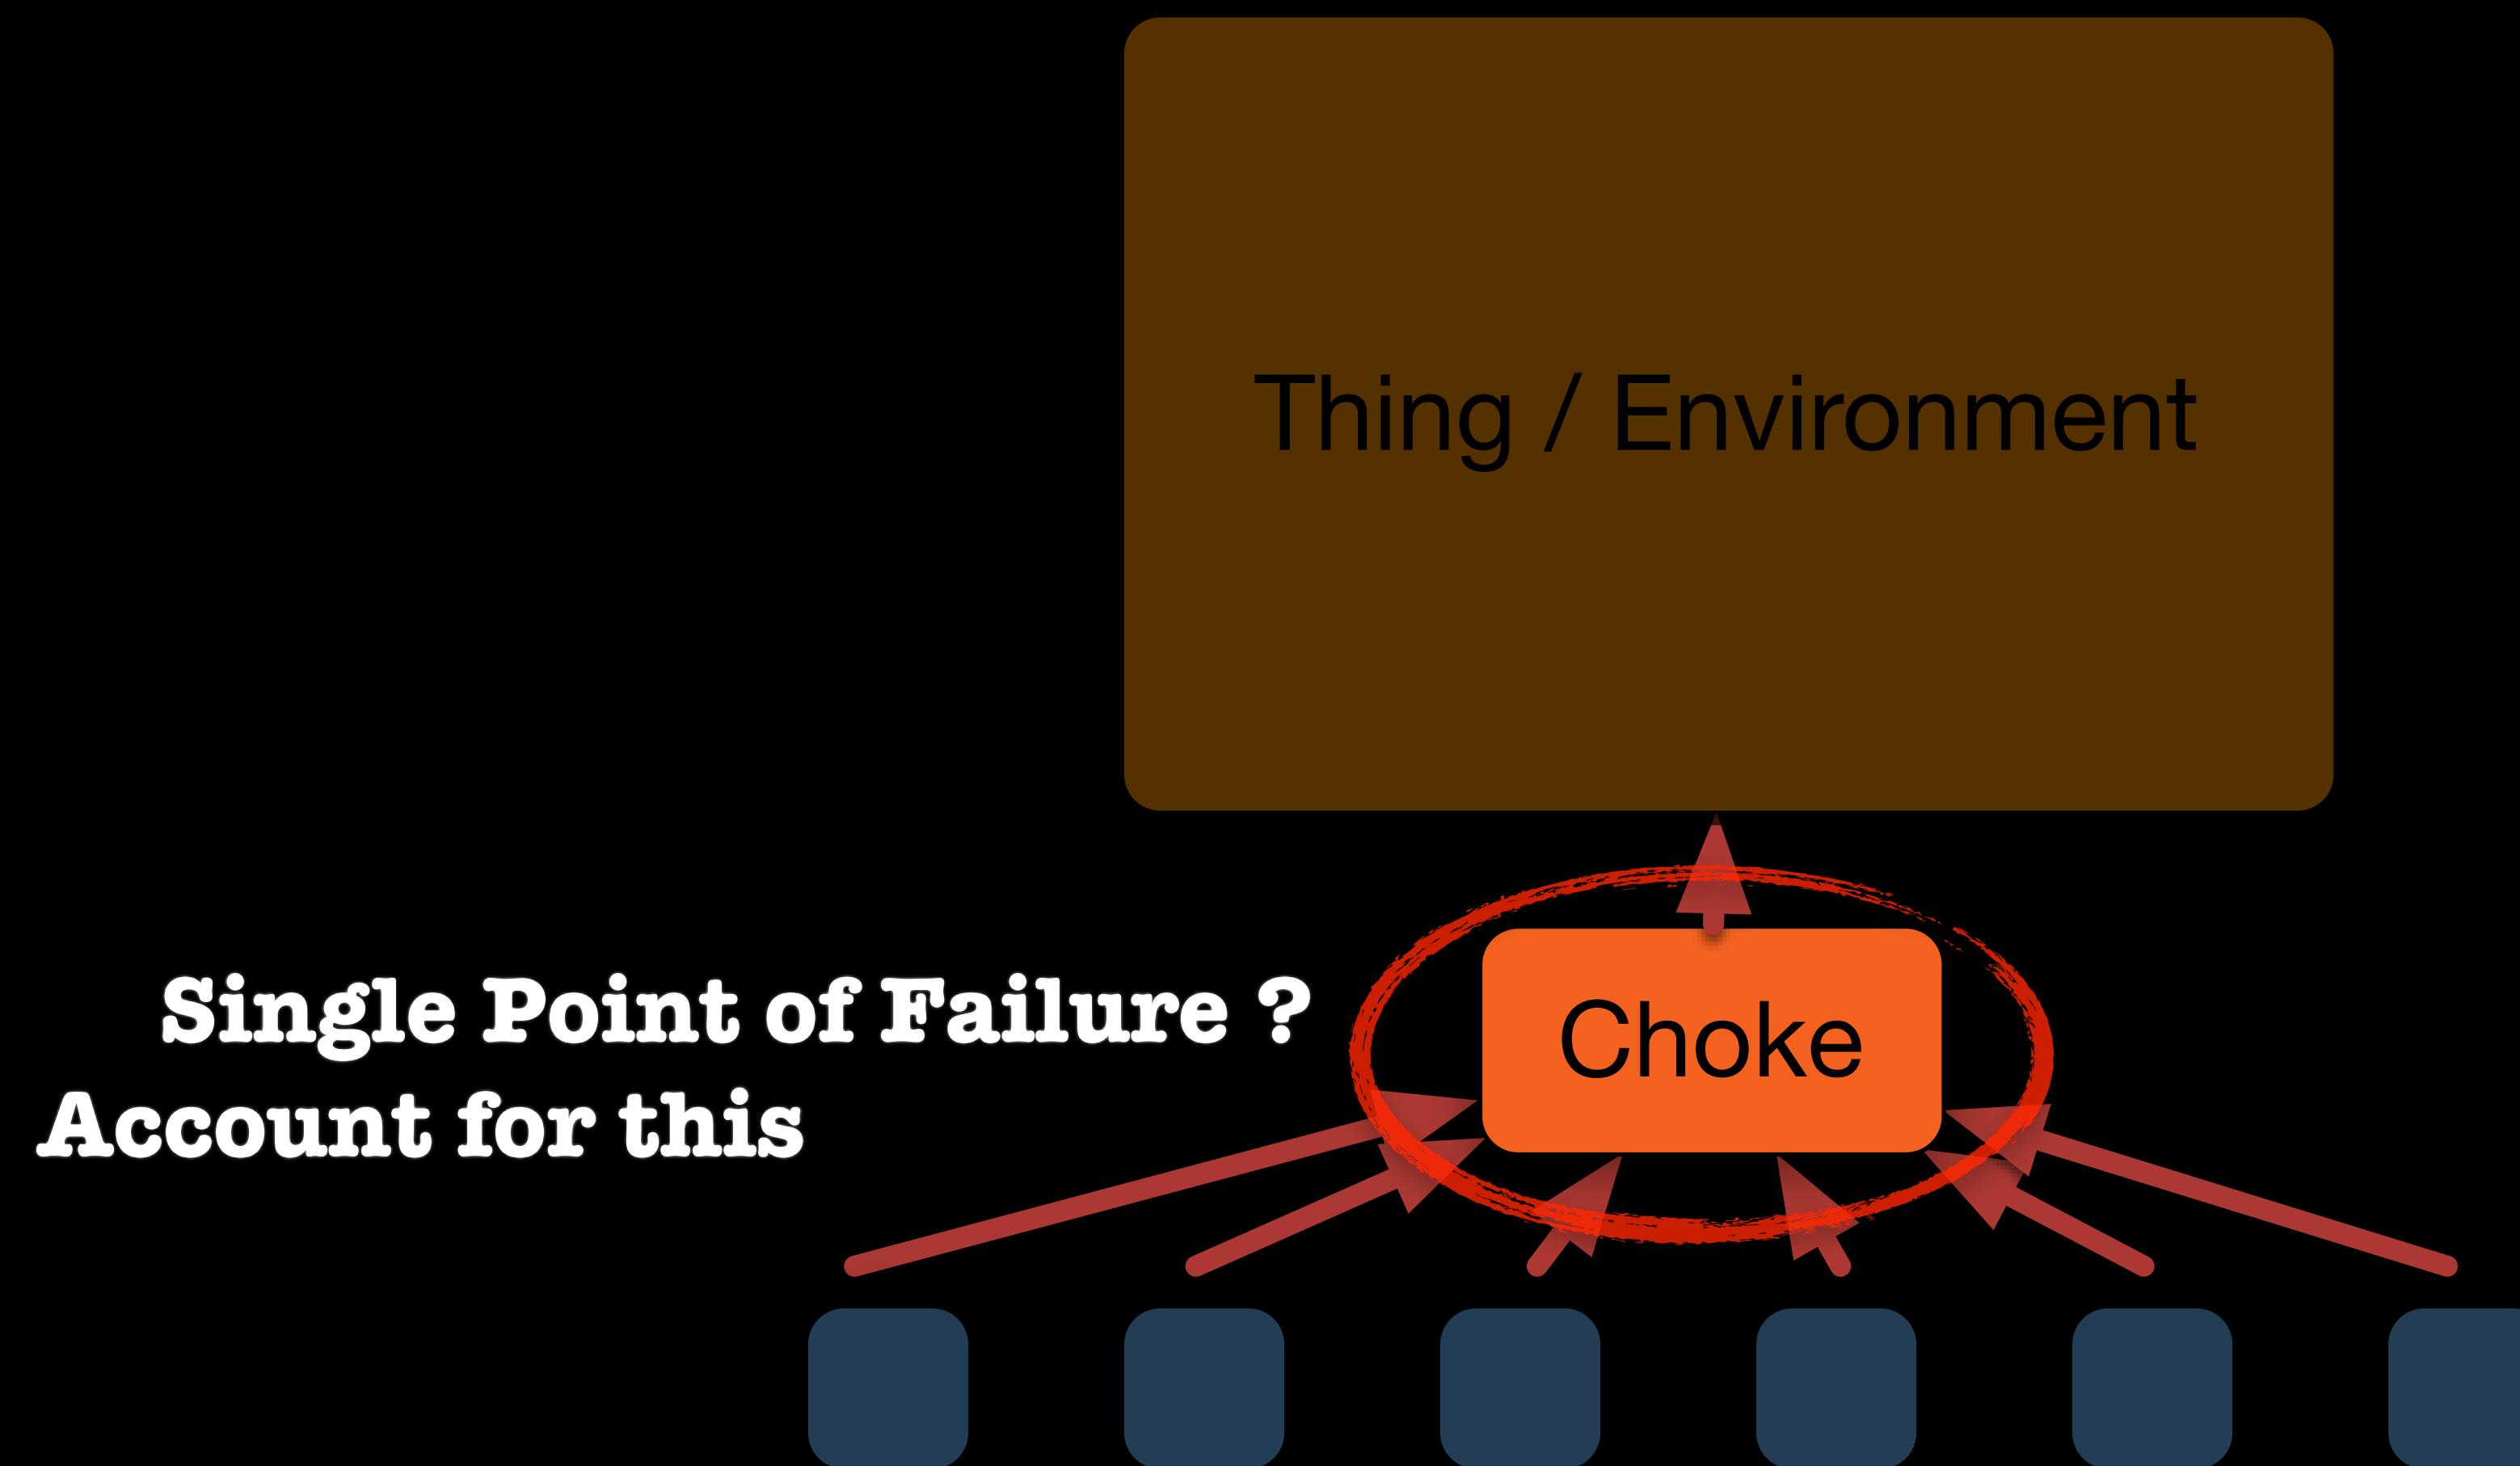
Architecting Choke Points



Architecting Choke Points



Architecting Choke Points



Assisting Sales



Sales and Customers

- ❖ Being part of Sales
- ❖ Customer Confidence
- ❖ Reducing involvement; scaling w/ sales
- ❖ Create a SIG (SharedAssessments.org)
- ❖ Create a Q&A/FAQ Repo
- ❖ Create training program for Sales Reps

Program Documentation

Reference Slide Only

- InfoSec Program and Policy Basis
- Team and Reporting Structure
- How you secure Assets
- How employees are trained (and background-checked)
- Processes Employee Use (logical controls)
- How Customer Data is Secured
- How systems are hardened and monitored
- Incident Response Plan Overview
- Security Communication and Notifications
- Development Practices / SDLC overview.
- Anything add'l relevant for regulatory reasons
- Docs / References (if applicable)

DevOps?

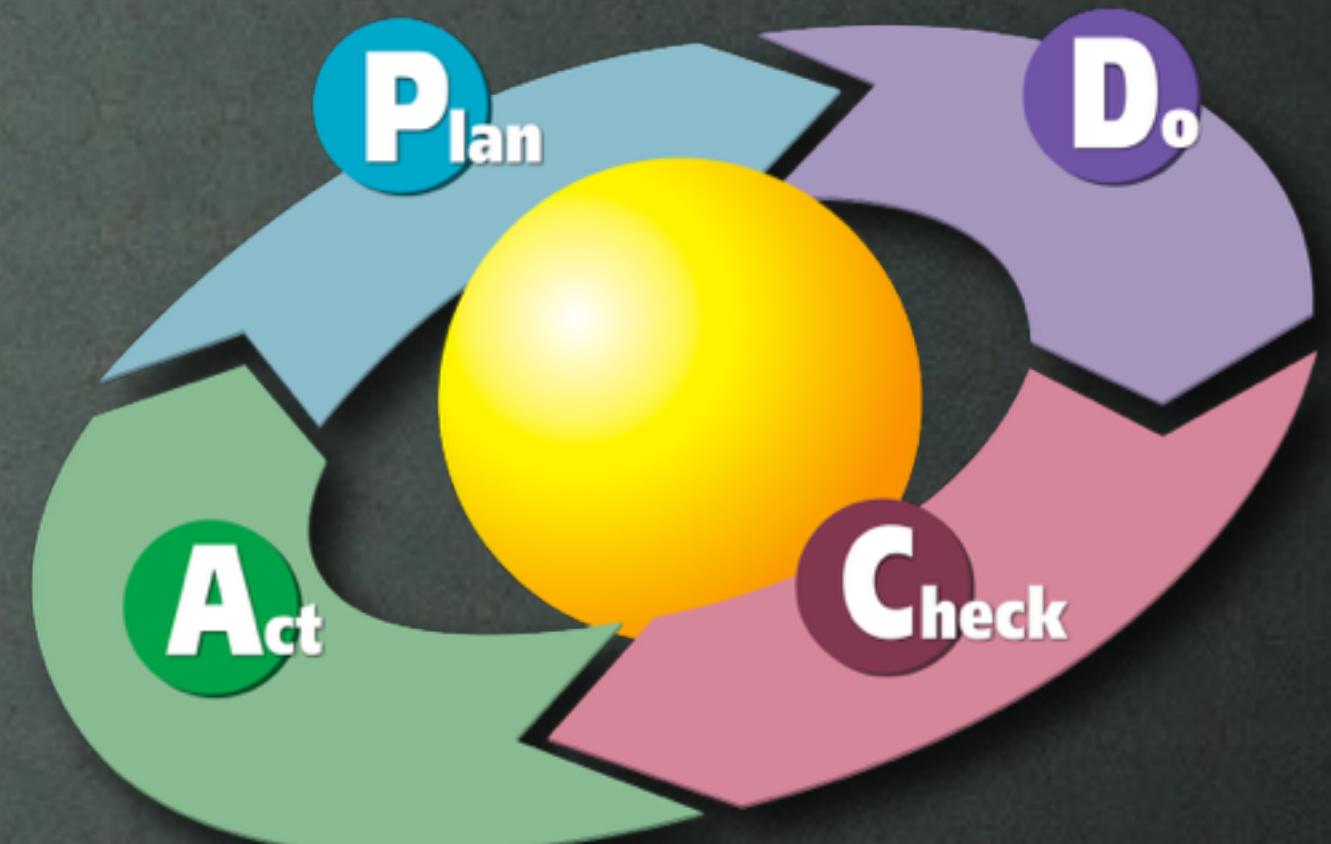


We're punting on this; go see
Building a Modern Security Engineering Organization - Zane Lackey

Auditing and QC

Are your automated processes working properly?

- ❖ 1. Manual. 2. Automated. 3. Verify/Improve. 4. Repeat.
- ❖ QA, Gap Analysis, API Updates
- ❖ Repeatability and Periodicity
- ❖ Docs Docs Docs
- ❖ Educate team(s) on workflow and process repeatability



AWS Specific

Reference Slide Only

- ❖ Use IAM Roles
- ❖ Enable CloudTrail (API logs) and EC2 Config (infra change logs) everywhere
- ❖ Use Multi Accounts (prod, dev/QA, logging-account, corp-intra, etc)
- ❖ Check blogs very frequently (weekly)
- ❖ Review Credential Report (ideally, automate)
- ❖ Tons of docs and videos (terminology can be confusing, and you'll hit information overload)
 - ❖ Require Tagging; kill things that don't follow your schema
 - ❖ Not all AWS best practices are realistic, recognize this.



Tools



Tools

- ❖ GitHub Commit Crawler
- ❖ S3 Key Auditor
- ❖ Nessus Process Controller
- ❖ AWS Enumeration Library
- ❖ HTTPS Uploader Tool
- ❖ Baseline Measuring Spreadsheet

Github Commit Crawler

[**https://github.com/jfalken/github_commit_crawler**](https://github.com/jfalken/github_commit_crawler)

- ❖ Monitoring for mistaken commits of secrets
- ❖ 1. Enumerate Github Org Members
- ❖ 2. Check all their recent commits
- ❖ 3. Rinse, Repeat
- ❖ Easy install and config; all packaged in a Docker Container

Github Commit Crawler

User	Audit_Date	Matched Keyword	File	Other Links
sr527	2015-04-01	password	auth_test.go	Diff Commit HTML
llvtt	2015-04-01	password	sharded_clusters.py	Diff Commit HTML
llvtt	2015-04-01	password	sharded_clusters.py	Diff Commit HTML
gabrielrussell	2015-04-01	password	nid.go	Diff Commit HTML
gabrielrussell	2015-04-01	secret	nid.go	Diff Commit HTML

- ❖ 3. Rinse, Repeat
- ❖ Easy install and config; all packaged in a Docker Container

Github Commit Crawler

User	Audit Date	Matched Keyword	File	Other Links
Found "secret" in https://api.github.com/repos/jmikola/arsjerm/commits/273308a0a265eecd6344d35ac548142e430874fb				
<pre>@@ -1,3 +1,5 @@ +# This file is a "template" of what your parameters.yml file should look like parameters: - locale: en - secret: ThisTokenIsNotSoSecretChangeIt + mongodb_uri: mongodb://127.0.0.1:27017 + locale: en + secret: ThisTokenIsNotSoSecretChangeIt</pre>				
User	Audit Date	Matched Keyword	File	Other Links

Github Commit Crawler

Monitoring found "secret" in https://github.com/gabrielrussell/5/commit/6344d35ac548142e430874fb

Enumeration found "secret" in https://github.com/gabrielrussell/5/commit/6344d35ac548142e430874fb

This file is a "template" of parameters:

- locale: en
- secret: ThisTokenIsNotSoSecret
- + mongodburi: mongodb://127.0.0.1:27017
- + locale: en
- + secret: ThisTokenIsNotSoSecret

Easy install and configuration.

Edit Configuration.

You currently don't have a config setup!

Edit your config file here.

GitHub Username: Your Github Username
AccessToken: Github Personal Access Token
Org_Name: The org you wish to audit

Hitting 'submit' will overwrite the existing config and restart the `ghcc` process. View the status on the [Supervisor](#) page.

Username:

AccessToken:

Org_Name:

Submit and Restart Crawler

Other Links

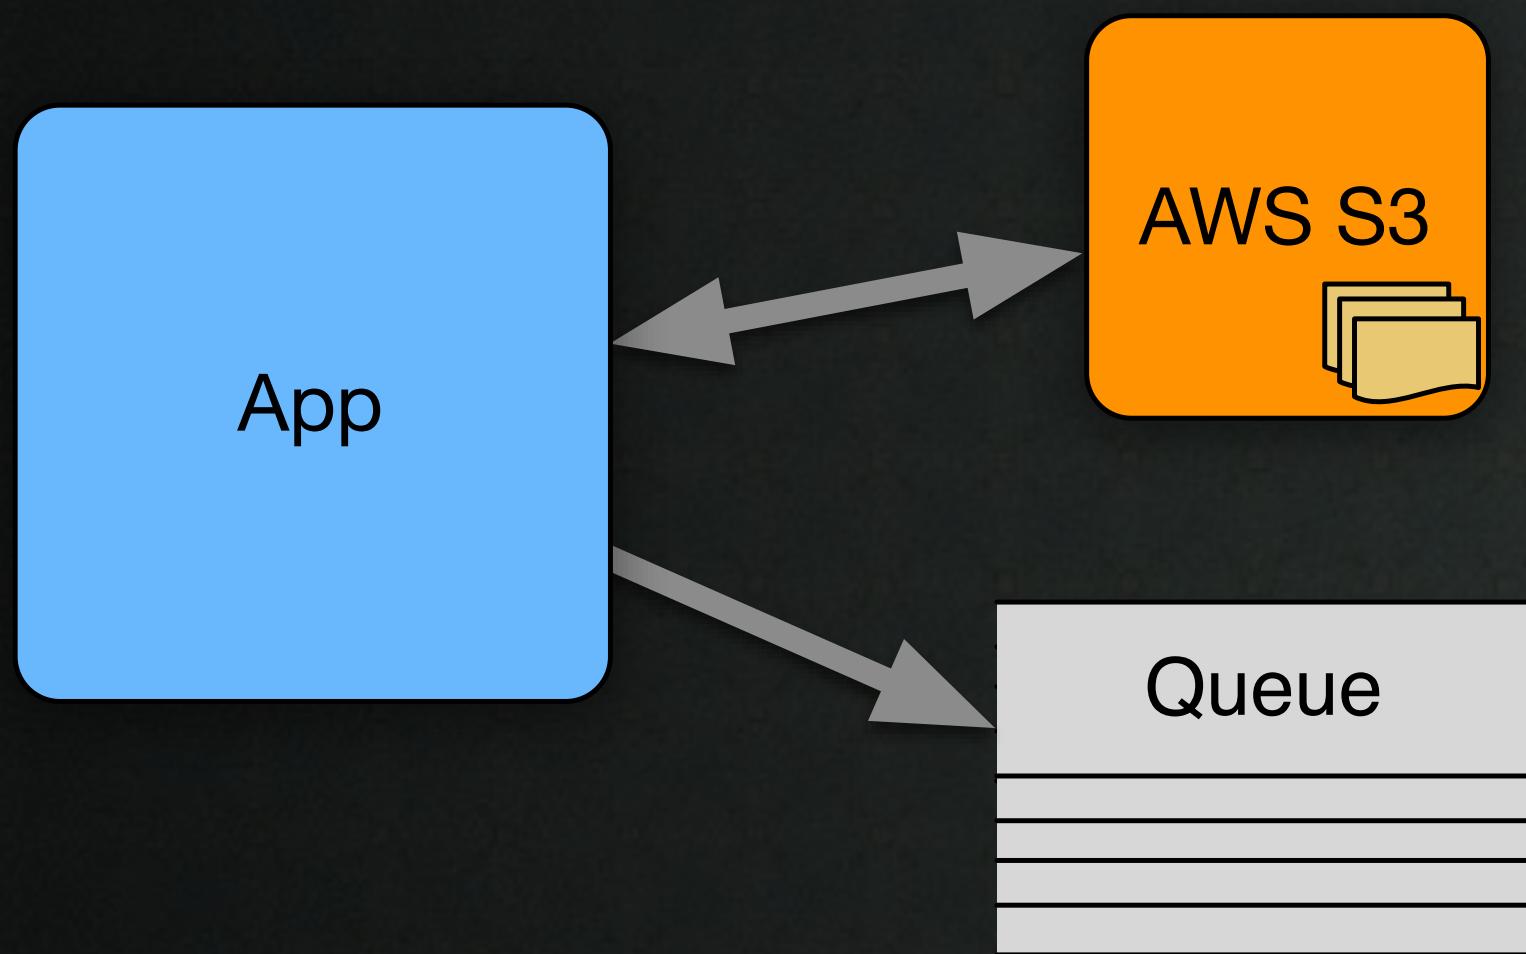
- 6344d35ac548142e430874fb Diff Commit HTML
- Diff Commit HTML

S3 Key Auditor

https://github.com/jfalken/s3_key_audit

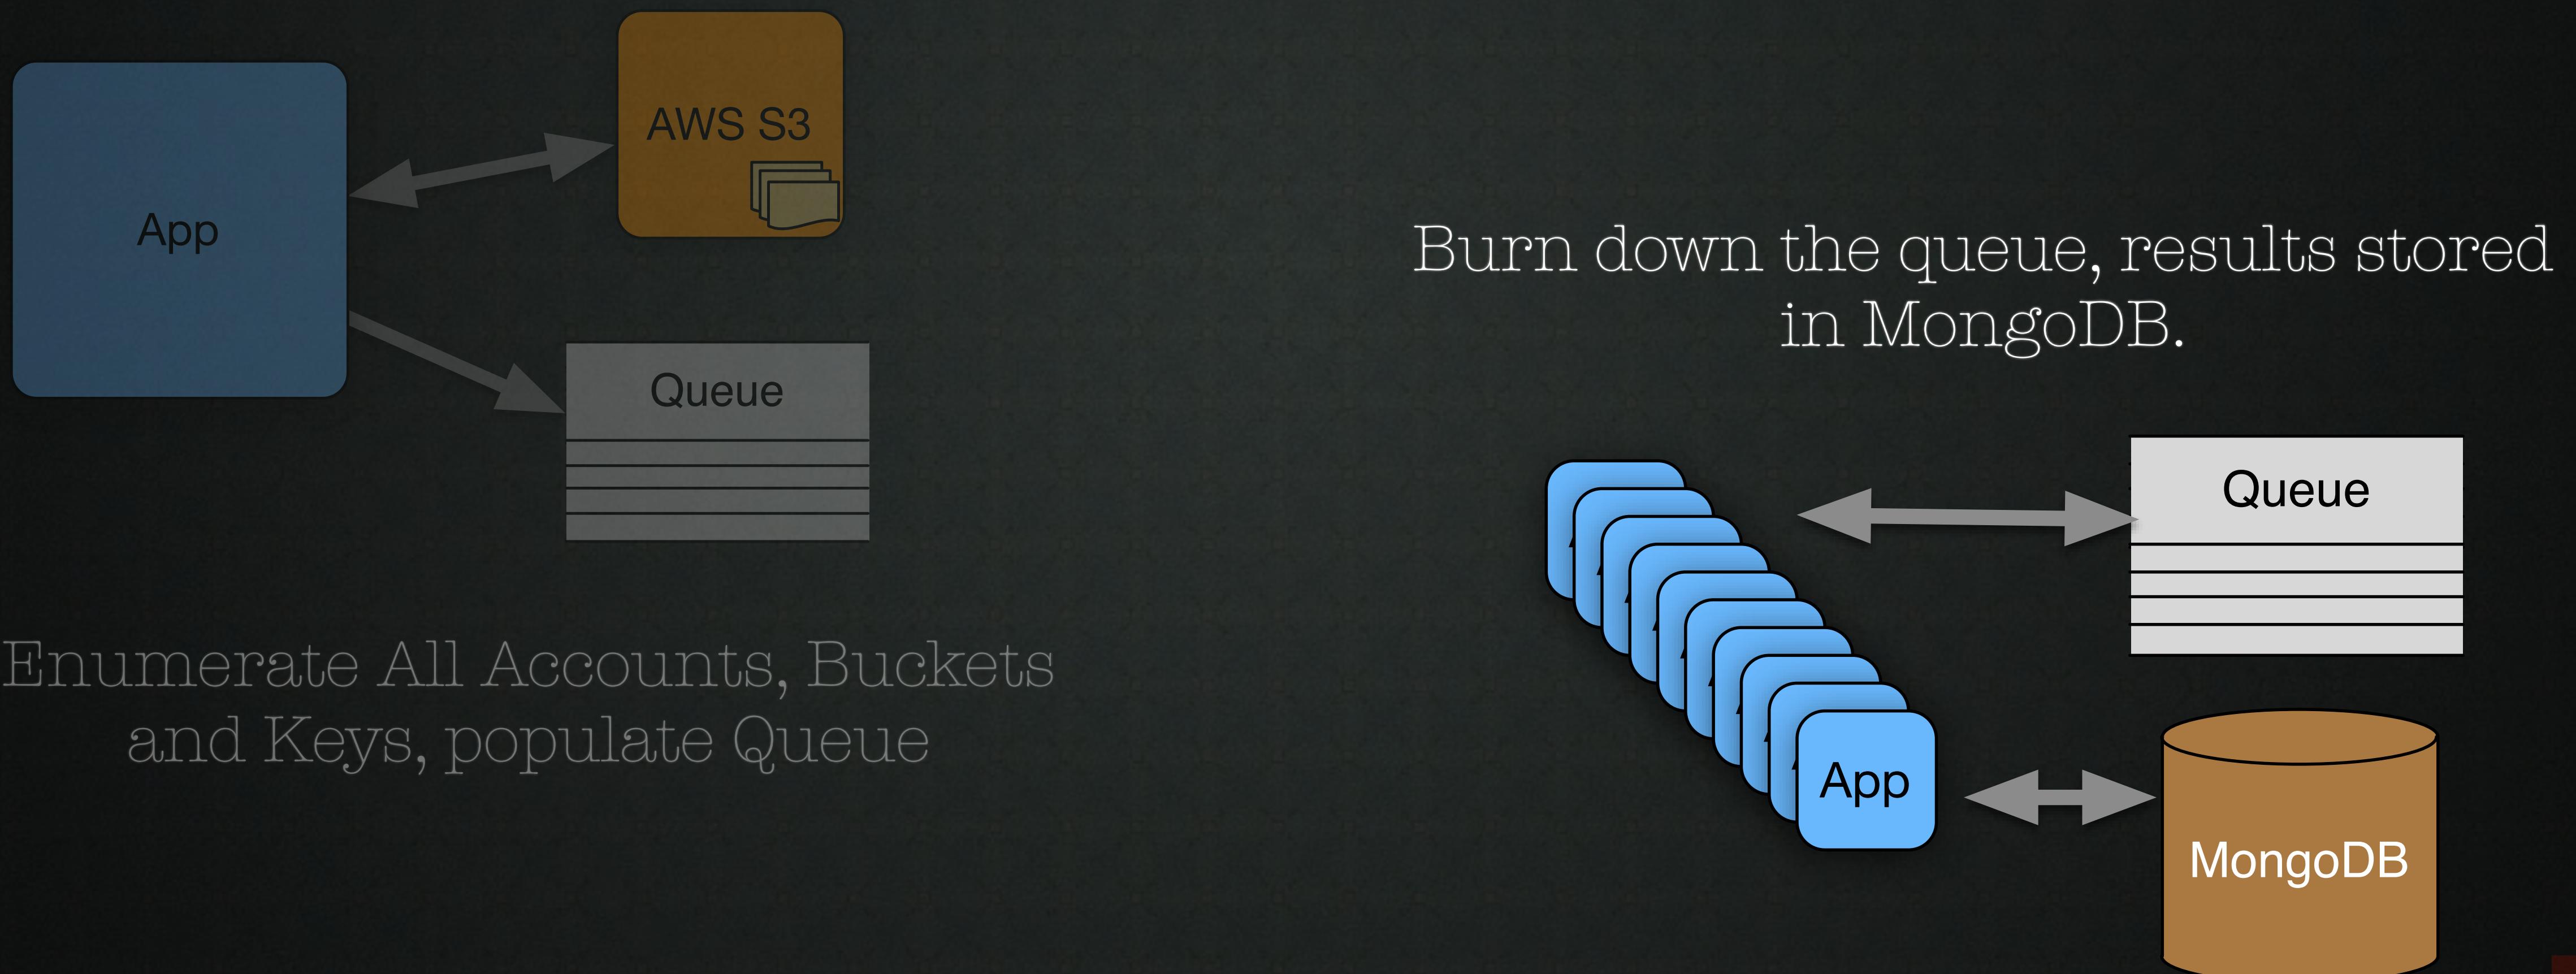
- ❖ S3 usage creeps over time; be mindful of what's going on
- ❖ ACLs and permissions on buckets don't matter, each key can have individual permissions
- ❖ Therefore; you need to audit every single key's permissions
- ❖ `s3_key_auditor.py`
 - ❖ Enumerate all keys in all buckets in all accounts
 - ❖ Multi-threaded app to burn through the queue as fast as possible
 - ❖ 2 Millions key ACL checks per hour (MBP 16GB RAM)

S3 Key Auditor

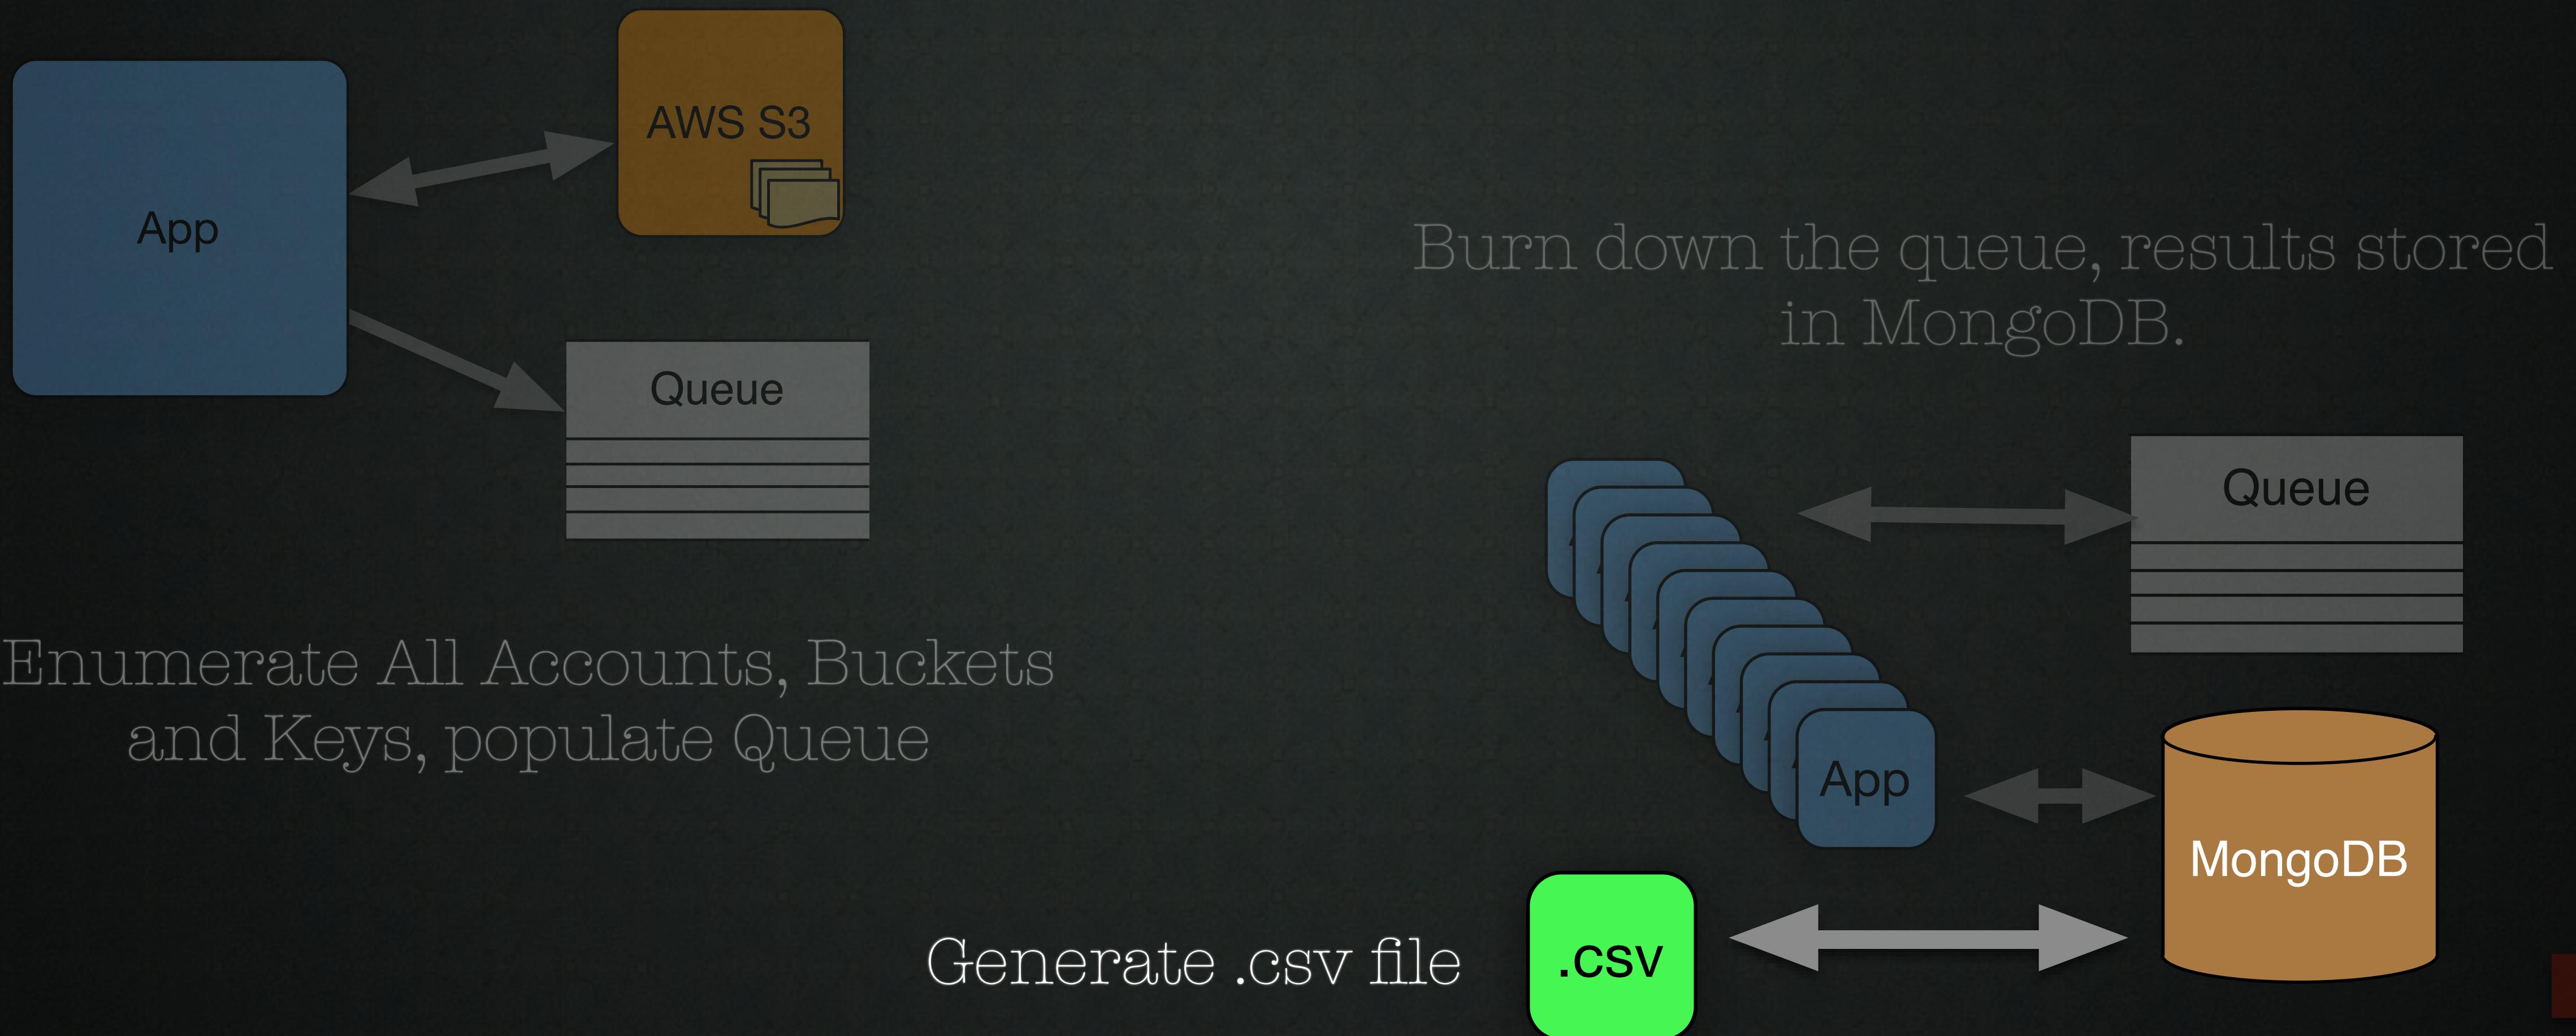


Enumerate All Accounts, Buckets
and Keys, populate Queue

S3 Key Auditor



S3 Key Auditor



Nessus Process Controller

https://github.com/jfalken/nessus_process_controller

- ❖ Enumerate, Launch Scans, Get Results, continuously.
- ❖ Supports AWS enumeration and also manually entered assets
- ❖ Scan as frequently as you like
- ❖ Get results in HTML, XML and JSON. Use output to drive other audit processes, cross-reference, etc.

AWS Enumeration Library

[**https://github.com/jfalken/aws_enumeration_lib**](https://github.com/jfalken/aws_enumeration_lib)

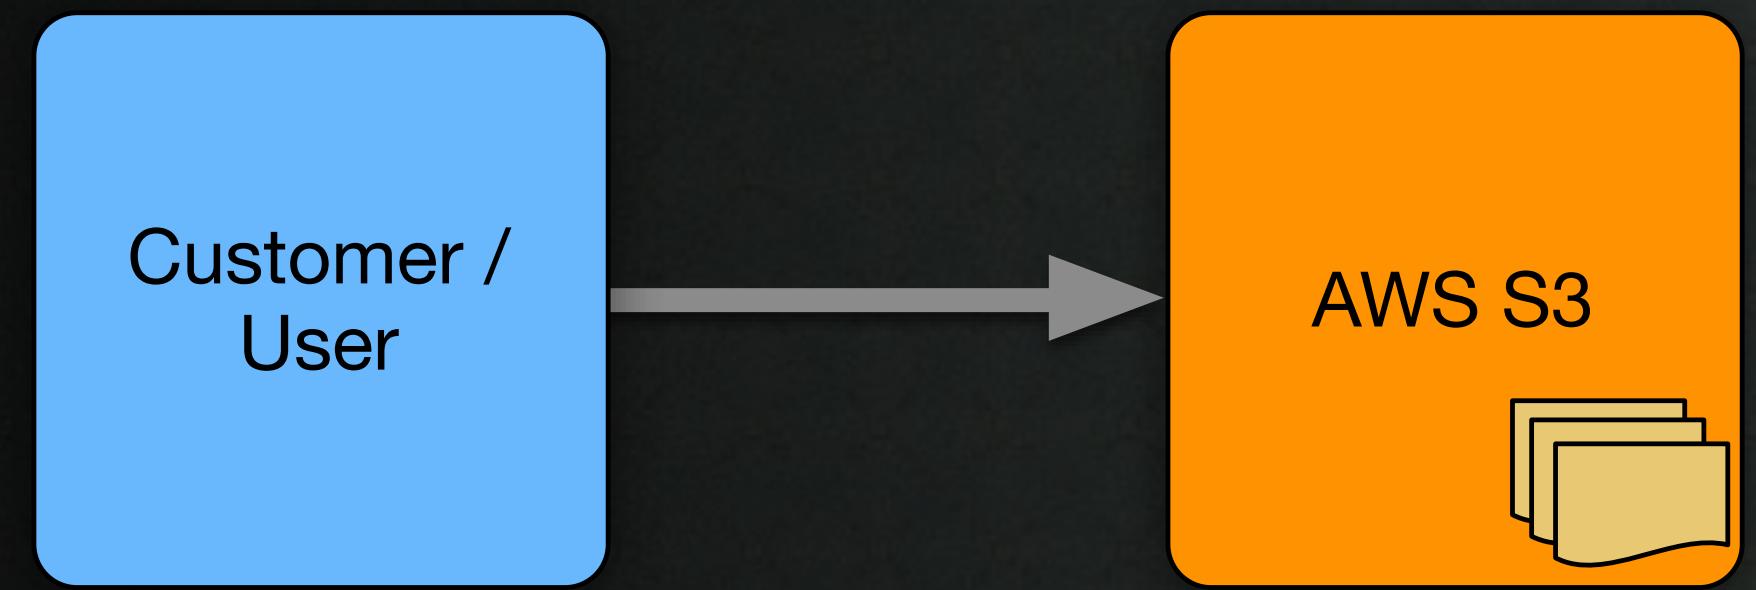
- ❖ We advocate multiple AWS Accounts
- ❖ Enumeration becomes burdensome
- ❖ AWSEnumerationObject
- ❖ Python library that wraps boto to make multi account and region asset identification easier

HTTPS Uploader

https://github.com/jfalken/s3_https_upload_portal

- ❖ Enable fast uploads to S3, with a minimal maintenance as possible.
- ❖ TLS, Encryption at Rest, Data Life Cycle Management
- ❖ Direct to S3 HTTPS Uploads, “GUI” drag-drop. curl support
- ❖ Thin web app to list/download file (employees don't need API creds)
- ❖ Py flask app on Elastic Beanstalk (ephemeral instance)

HTTPS Uploader



Customer /
User

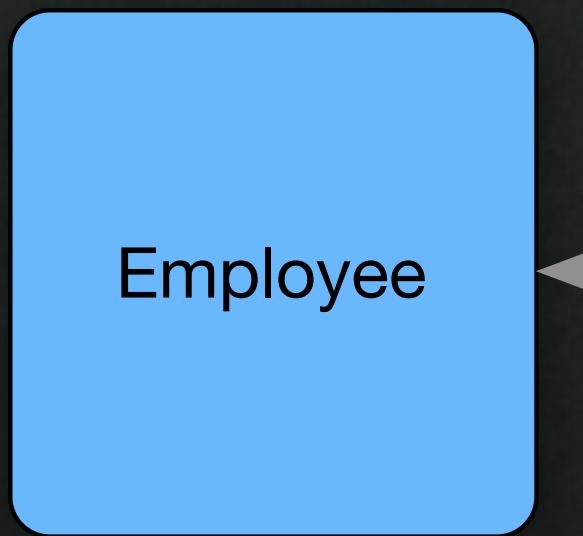
AWS S3

Customer Uploads (presigned URLs)

TLS, Encrypt At Rest, Direct to S3

Employee Access

Enable access for employees w/o S3 Creds



Employee

Python Flask /
Elastic
Beanstalk

AWS S3

HTTPS Uploader

Baseline Spreadsheet

- ❖ Based on ISO 27002:2015 Controls
- ❖ Loosely establish questions based on the control domains
- ❖ Scope yourself, question yourself, do this regularly
- ❖ Reflect on the output; wheres your lowest point? Make progress on that next quarter
- ❖ REPO: TBD

Moving
Forward ...

Open Source ‘Book’



The content of these slides will be made into a ‘book’ format, open sourced on GitHub for collaboration. End goal: to help others starting out in this area.

References

- ❖ Give Credit Where its due
- ❖ Check out their talks and work
 - ❖ Zane Lackey
 - ❖ Rich Smith
 - ❖ Haroon Meer
 - ❖ Chris Evans
 - ❖ Dan Geer
 - ❖ Chris Rohlf

end